

Time Series: A First Course with Bootstrap Starter

Contents

Lesson 3: Takes and tools	3
Takes	3
Questions	3
Links to JD+	3
Overview in book	3
Concepts	3
R skills	3
Lesson 3-1: Nonparametric Smoothing	3
My add-ons	3
Tucker questions	3
TP / questions	3
Insert, flesh out	3
Concepts	3
Code skills	3
Nonparametric Regression	4
Kernel and Bandwidth	4
Edge Effects	4
Nonparametric Regression for Population	4
Lesson 3-2: Linear Filters	5
Example 3.2.5. Simple Moving Average	6
Linear Time Series	6
Smoothing of Population	6
Lesson 3-3: Examples of Filters	7
Paradigm 3.3.1. Smoother	7
Example 3.3.2. Smoothers Applied to Gasoline Sales	7
Paradigm 3.3.3. Difference Filter	240
Example 3.3.4. Differencing Applied to Gasoline Sales	240
The Backward Shift	241
Powers of Backward Shift	241
Simple Moving Average Filter	241
Differencing Filter	241
Lesson 3-4: Trends	242
Example 3.4.1. Western Housing Starts Cycle	242
Paradigm 3.4.12. Nonparametric Trend Estimation	244
Paradigm 3.4.15. Trend Elimination	244
Lesson 3-5: Seasonality	245
Example 3.5.1. Mauna Loa Growth Rate	246
Seasonal Regression	247
Seasonal Moving Average	248
Example 3.5.8. Seasonal Moving Average for Motor Retail Sales	248

Remark 3.5.9. Seasonal Adjustment	249
Example 3.5.10. Seasonal Adjustment of Mauna Loa Growth Rate	249
Lesson 3-6: Trend and Seasonality Together	250
Proposition 3.6.5	250
Example 3.6.7. Trend Filters that Eliminate Seasonality	250
Paradigm 3.6.8. Classical Decomposition from Trend Filtering	251
Example 3.6.13. Classical Decomposition of Western Housing Starts via Linear Filtering	251
Lesson 3-7: Integrated Processes	252
Example 3.7.1. Random Walk	252
Integrated Process	253
Example of Integrated Moving Average	253
Example 3.7.3. Unit Root Process	254

Lesson 3: Takes and tools

Takes

Questions

key = small TPs on my data

Links to JD+

Overview in book

Concepts

R skills

Lesson 3-1: Nonparametric Smoothing

- The filtering of time series uses some concepts from nonparametric smoothing.

My add-ons

Tucker questions

- around hp filter

TP / questions

- apply to RF3030
- regression with dummies (equivalent to trig regression)
- numeric effect of seasonal aggregation filter
- why 1-Bs not directly applied in x11: removes T and S nothing left
- what we want is trend filters that eliminate seasonality = 2×12 classical trend estimate
- recréer la 2×12 to prove it eliminates stable S (diff de M1+M2)

Insert, flesh out

- differencing filters from B4
- HP filter
- here code giving an HP filter for seasonal data

Concepts

- convolution (see P Course)
- Laurent Series (2 side power series)

Code skills

- ma filter from first principles (att NAs)
- ma filter with stats filter function (convolution)
- movie !
- compSmooth function (x11 surrogate)
- hp filter revu ?
- simulate random walk as cummulation of iid or MA(1)

Options Set `xaxt = "n"` and `yaxt = "n"` to remove the tick labels of the plot

Nonparametric Regression

- We might think of the time series mean $\mu_t = \mathbb{E}[X_t]$ as an arbitrary, smoothly varying function.
- A very simple case:

$$X_t = \mu_t + Z_t$$

with $Z_t \sim i.i.d.(0, \sigma^2)$.

Here X_t not stationnary, Z_t iid quite a big assumption

- We may estimate μ_t via averaging over neighboring values, e.g.,

$$\hat{\mu}_t = \frac{1}{2m+1} \sum_{k=-m}^m X_{t-k}.$$

This works because

$$\hat{\mu}_t = \frac{1}{2m+1} \sum_{k=-m}^m \mu_{t-k} + \frac{1}{2m+1} \sum_{k=-m}^m Z_{t-k},$$

and the first term is approximately equal to μ_t if m is small and the mean is smooth (not changing too much over time). Also, the second term has mean zero and variance $\sigma^2/(2m+1)$, which tends to zero as m increases to ∞ .

- So bias is lower for small m , but variance is lower for large m , creating a tension.

Kernel and Bandwidth

Bandwidth: half filter length, if smaller keeps more local behaviour

- Above, we have equal weights on the observations. We can have unequal weights, through the device of a *kernel*, which is a function that weights the data.
- The *bandwidth* m defines a neighborhood of values around time t , for which we use the kernel's weights.
- Large bandwidth yields more smoothing, with suppression of local features.
- Small bandwidth yields less smoothing, with higher variability.

Edge Effects

- We can only compute $\hat{\mu}_t$ for $t = m+1, \dots, n-m$. So the values for $t = 1, \dots, m$ and $t = n-m+1, \dots, n$ are “missing”.
- Sometimes this so-called “edge effect” is addressed by designing asymmetric smoothing at the boundaries.
- Often the edge effects are ignored, when we are interested in the interior of the time series sample.


Nonparametric Regression for Population

- Apply the nonparametric regression technique with $m = 20$ to U.S. population data.
- First, we write a function to do local averaging.

```
simple.ma <- function(x,m)
{
  weights <- rep(1,2*m+1)/(2*m+1) # div done elem by elem
  n <- length(x)
  trend <- NULL # init values of trend
  for(t in (m+1):(n-m)) # for the values which can be computed (n-m-1 ?)
  {
    trend <- c(trend,sum(x[(t-m):(t+m)]*weights))
  }
  trend <- c(rep(NA,m),trend,rep(NA,m))
  return(trend)
}
```

- Then we apply this to the data.

```
pop <- read.table("USpop.dat")
pop <- ts(pop, start = 1901)
m <- 30
pop.trend <- simple.ma(pop,m)
plot(pop,xlab="Year",ylab="U.S. Population",col=1,lwd=1,yaxt="n",xaxt="n")
lines(ts(pop.trend,start=1901,frequency=1),lty=1,lwd=1,col=2)
```



Lesson 3-2: Linear Filters

- A *linear filter* maps an input time series $\{X_t\}$ to an output time series $\{Y_t\}$ by taking a linear combination of past, present, and future observations:

$$Y_t = \sum_{k \in \mathbb{Z}} \psi_k X_{t-k}.$$

The *filter coefficients* (or *filter weights*) are $\{\psi_k\}$. Note the convention that ψ_k weights an observation occurring k time points in the past (viewing time t as the present).

Written as a convolution.

Example 3.2.5. Simple Moving Average

- Recall from nonparametric regression, we may estimate a slowly-changing mean via averaging over neighboring values:

$$\frac{1}{2m+1} \sum_{k=-m}^m X_{t-k}.$$

This weights the past m and the future m observations equally, and is an example of a linear filter. It also called a *moving average*, because the observations are averaged (over a window of $2m+1$ time points) in a way that moves over the time series. A *simple moving average* has equal weights (a general moving average could have unequal weights).

Linear Time Series

- A time series $\{Y_t\}$ is said to be *linear* if it is defined as the output of a linear filter (with coefficients $\{\psi_k\}$) applied to $\{X_t\}$, and $X_t \sim i.i.d.(\mu, \sigma^2)$.
- The mean of Y_t is

$$\mathbb{E}[Y_t] = \mu \sum_{k \in \mathbb{Z}} \psi_k.$$

- The autocovariance of $\{Y_t\}$ is

$$\text{Cov}[Y_t, Y_{t-h}] = \sum_{j,k \in \mathbb{Z}} \psi_j \psi_k \text{Cov}[X_{t-j}, X_{t-h-k}] = \sigma^2 \sum_{k \in \mathbb{Z}} \psi_k \psi_{k+h}.$$

This calculation is obtained by seeing that $\text{Cov}[X_{t-j}, X_{t-h-k}]$ is zero unless $j = k + h$, in which case it equals σ^2 .

Linear times series are linear combination of WN More generally if input statyionary, output as well.

Smoothing of Population


- We apply a simple moving average to U.S. Population, using the R *filter* function.
- Overlay different choices of m .

```
pop <- read.table("USpop.dat")
pop <- ts(pop, start = 1901)

m <- 20
ma.weights <- rep(1,2*m+1)/(2*m+1)
pop.trend1 <- stats::filter(pop,ma.weights,method="convolution",sides=2) # function from stats

m <- 5
ma.weights <- rep(1,2*m+1)/(2*m+1)
pop.trend2 <- stats::filter(pop,ma.weights,method="convolution",sides=2)

plot(pop,xlab="Year",ylab="U.S. Population",col=1,lwd=1,yaxt="n",xaxt="n")
lines(ts(pop.trend1,start=1901,frequency=1),lty=1,lwd=1,col=2)
lines(ts(pop.trend2,start=1901,frequency=1),lty=1,lwd=1,col=3)
```



Lesson 3-3: Examples of Filters

- We explore smoothers, differencing, and the backshift operator.

Paradigm 3.3.1. Smoother

- Linear filters that suppress oscillations and reveal long-term trends are called *smoothers*.
- A simple moving average is an example of a smoother.


Low pass filters Used for trend estimation

Example 3.3.2. Smoothers Applied to Gasoline Sales

- We apply a simple moving average with $m = 10$ to the logged seasonally adjusted gasoline series.


```
gassa <- read.table("GasSA_2-11-13.dat")
gassa.log <- ts(log(gassa),start=1992,frequency=12)


h <- 10
simple.ma <- rep(1,2*h+1)/(2*h+1) # coeffs
gas.trend <- stats::filter(gassa.log,simple.ma,method="convolution",sides=2)
gas.trend <- ts(gas.trend,start=1992,frequency=12)
plot(ts(gassa.log,start=1992,frequency=12),col=1,ylab="",xlab="Year",lwd=2)
lines(ts(gas.trend,start=1992,frequency=12),col=2,lwd=2)
```





- We can run a visualization in the notebook.


```
### Movie
movie <- TRUE
delay <- 0
flex <- 0
siglen <- length(gas.trend) - 2*h
range <- seq(1,siglen,1)
if(movie) {
  for(t in range)
  {
    Sys.sleep(delay)
    filterweight <- c(rep(NA,t-1),simple.ma + gas.trend[h+t] + flex,
                      rep(NA,siglen-t))
    plot(ts(gassa.log,start=1992,frequency=12),col=1,ylab="",xlab="Year",lwd=2)
    lines(ts(gas.trend[1:(h+t)],start=1992,frequency=12),col=2,lwd=2)
    lines(ts(filterweight,start=1992,frequency=12),col="#00FF0090",lwd=2)
  }
}
```





















































































































































































































































































































































































































































































































































































































































































































Paradigm 3.3.3. Difference Filter

- To suppress long-term dynamics, we can *difference*:

$$Y_t = X_t - X_{t-1}.$$


This is a filter with weights $\psi_0 = 1$, $\psi_1 = -1$, and zero otherwise. It is called the *differencing filter*. This filter reduces polynomials in time by one degree; lines are reduced to constants.

One sided filter

Example 3.3.4. Differencing Applied to Gasoline Sales

- We apply differencing to logged seasonally adjusted gasoline sales. The result can be interpreted as a *growth rate*. (DL de $\log(1+x) = x$)

```
gas.diff <- diff(gassa.log)
par(oma=c(2,0,0,0),mar=c(2,4,2,2)+0.1,mfrow=c(2,1),cex.lab=.8)
plot(gassa.log,xlab="",ylab="Log Gas Sales",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
plot(gas.diff,xlab="",ylab = "Gas Growth Rate",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
mtext(text="Year",side=1,line=1,outer=TRUE)
```

The Backward Shift

- We define the *backward shift* (or *lag*) operator B , which shifts a time series back one time unit. This is a linear filter that lags time by one unit:

$$Y_t = X_{t-1}.$$

So the filter weights are $\psi_1 = 1$ and zero otherwise. Informally we write $BX_t = X_{t-1}$.

Powers of Backward Shift

- Then $X_{t-k} = B^k X_t$, and a general filter is expressed as

$$Y_t = \sum_{k \in \mathbb{Z}} \psi_k X_{t-k} = \sum_{k \in \mathbb{Z}} \psi_k B^k X_t,$$

so we write the filter as $\sum_{k \in \mathbb{Z}} \psi_k B^k$. Call this $\Psi(B)$. (Mathematically, $\Psi(z)$ is a Laurent series, as a usual power series is one sided and here it is two sided)

Simple Moving Average Filter

- Expressed as $\Psi(B) = (2m+1)^{-1} \sum_{k=-m}^m B^k$.

Differencing Filter

- Expressed as $\Psi(B) = 1 - B$.

Lesson 3-4: Trends

- Trends measure the long-term behavior of a time series.
- Trends are steady movements in the level of a time series.
- We can use smoothers to estimate trends.
- Removing a trend by subtraction allows us to see other features.

Example 3.4.1. Western Housing Starts Cycle

- We estimate and eliminate the trend in the Westing Housing Starts data, so that we can more clearly see the business cycle movements.

```
hpsa <- function(n,period,q,r)
{
  # hpsa
  #   gives an HP filter for seasonal data
  #   presumes trend+seas+irreg structure
  #   trend is integrated rw
  #   seas is seasonal rw
  #   irreg is wn
  #   q is snr for trend to irreg
  #   r is snr for seas to irreg

  # define trend differencing matrix

  delta.mat <- diag(n)
  temp.mat <- 0*diag(n)
  temp.mat[-1,-n] <- -2*diag(n-1)
  delta.mat <- delta.mat + temp.mat
  temp.mat <- 0*diag(n)
  temp.mat[c(-1,-2),c(-n,-n+1)] <- 1*diag(n-2)
  delta.mat <- delta.mat + temp.mat
  diff.mat <- delta.mat[3:n,]

  # define seasonal differencing matrix

  delta.mat <- diag(n)
  temp.mat <- 0*diag(n)
  inds <- 0
  for(t in 1:(period-1))
  {
    temp.mat <- 0*diag(n)
    temp.mat[-(1+inds),-(n-inds)] <- 1*diag(n-t)
    delta.mat <- delta.mat + temp.mat
    inds <- c(inds,t)
  }
  sum.mat <- delta.mat[period:n,]

  # define two-comp sig ex matrices

  #trend.mat <- solve(diag(n) + t(diff.mat) %*% diff.mat/q)
  #seas.mat <- solve(diag(n) + t(sum.mat) %*% sum.mat/r)
  trend.mat <- diag(n) - t(diff.mat) %*% solve(q*diag(n-2) + diff.mat %*%
    t(diff.mat)) %*% diff.mat
  seas.mat <- diag(n) - t(sum.mat) %*% solve(r*diag(n-period+1) + sum.mat %*%
```

```

t(sum.mat)) %*% sum.mat

# define three-comp sig ex matrices

trend.filter <- solve(diag(n) - trend.mat %*% seas.mat) %*%
  trend.mat %*% (diag(n) - seas.mat)
seas.filter <- solve(diag(n) - seas.mat %*% trend.mat) %*%
  seas.mat %*% (diag(n) - trend.mat)
irreg.filter <- diag(n) - (trend.filter + seas.filter)

filters <- list(trend.filter,seas.filter,irreg.filter)
return(filters)
}


Wstarts <- read.table("Wstarts.b1",skip=2)[,2]
Wstarts <- ts(Wstarts,start = 1964,frequency=12)
n <- length(Wstarts)
q <- .0001
r <- 1
hp.filters <- hpsa(n,12,q,r)

wstarts.trend <- hp.filters[[1]] %*% Wstarts
wstarts.seas <- hp.filters[[2]] %*% Wstarts
wstarts.cycle <- hp.filters[[3]] %*% Wstarts
wstarts.sa <- wstarts.trend + wstarts.cycle

comps <- ts(cbind(wstarts.trend,wstarts.seas,wstarts.cycle),start=1964,frequency=12)
trend <- ts(wstarts.trend,start=1964,frequency=12)
cycle <- ts(wstarts.cycle,start=1964,frequency=12)

par(oma=c(2,0,0,0),mar=c(2,4,2,2)+0.1,mfrow=c(3,1),cex.lab=.8)
plot(Wstarts, ylab="West Starts",xlab="",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
plot(trend,xlab="",ylab = "Trend",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
plot(cycle,xlab="",ylab = "Business Cycle",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
mtext(text="Year",side=1,line=1,outer=TRUE)

```



Paradigm 3.4.12. Nonparametric Trend Estimation

- We consider trends μ_t to be either a deterministic (but unknown) function of time, or are an unobserved stochastic process:

$$X_t = \mu_t + Z_t.$$

Here $\mathbb{E}[X_t] = \mu_t$ and $\{Z_t\}$ is a mean zero time series.

- A *two-sided moving average* is a linear filter that can be used to estimate trends:

$$\Psi(B) = \sum_{k=-m}^m \psi_k B^k.$$

When all the weights are equal, this is a simple moving average. It is two-sided because the moving average uses past and future data:

$$\hat{\mu}_t = \Psi(B)X_t = \sum_{k=-m}^m \psi_k X_{t-k}.$$

- A *symmetric moving average* weights future and past equally, i.e., $\psi_k = \psi_{-k}$.
- As an estimator $\hat{\mu}_t$ of the unknown trend μ_t , there is less bias if m is small, and less variance if m is large. This is a **Bias-Variance Dilemma**.

Paradigm 3.4.15. Trend Elimination

2 ways to eliminate trend (very different results)


- We can remove a trend by subtracting off its estimate.

- We can also apply the differencing filter.
- Example with US population, detrended by a $m = 10$ simple moving average and by the difference filter.

```
pop <- read.table("USpop.dat")
pop <- ts(pop, start = 1901)

p <- 10
pop.smooth <- stats::filter(pop, rep(1, 2*p+1)/(2*p+1), method="convolution")
pop.smooth <- ts(pop.smooth, start= 1901)
pop.diff <- diff(pop)

par(mar=c(4,4,2,2)+0.1, cex.lab=.8)
plot(ts(pop - pop.smooth, start=1901), lwd=1, lty=3, ylim=c(-3.5*10^6, 3.5*10^6),
     ylab="U.S. Population", xlab="Year", yaxt="n", xaxt="n")
axis(1, cex.axis=.5)
axis(2, cex.axis=.5)
lines(pop.diff, col=gray(.8), lwd=2)
```




Lesson 3-5: Seasonality

- Regular patterns of a periodic nature are called *seasonality*.
- The same notion as a *cycle*, but is associated with the calendar.
- We may wish to estimate and remove seasonality in order to see other patterns.

Example 3.5.1. Mauna Loa Growth Rate


- We consider the Mauna Loa CO2 time series.

```
mau <- read.table("mauna.dat",header=TRUE,sep="")  
mau <- ts(mau,start=1958,frequency=12)  
plot(mau,xlab="Year",ylab="Mauna Loa")
```



- We can difference the series to eliminate the trend.
- A pronounced seasonal pattern is apparent; but it is not completely periodic.


```
mau.diff <- diff(mau)  
plot(mau.diff,xlab="Year",ylab="Mauna Loa")
```



Seasonal Regression

- An exactly periodic effect would be $X_t = X_{t-s}$ for seasonal period s (e.g., $s = 12$ for monthly time series), and any t .
- If this were true, we could use dummy regressors. Usually this works poorly; if the series were periodic, then the filter $1 - B^s$ would annihilate it. VERY good point here
- Example of Mauna Loa; seasonal differencing of the growth rate leaves a non-zero time series.

```
mau.diffs <- diff(mau.diff,lag=12)
plot(mau.diffs,xlab="Year",ylab="Mauna Loa")
```



Seasonal Moving Average

- A better way is to estimate seasonality using a *seasonal moving average*. This is a moving average filter where only coefficients with indices that are a multiple of s are non-zero.
- This means

$$\Psi(B) = \sum_{k=-ms}^{ms} \psi_k B^k = \sum_{j=-m}^m \psi_{js} B^{js}$$

is a seasonal moving average, because $\psi_k = 0$ unless $k = js$ for some integer j .

Example 3.5.8. Seasonal Moving Average for Motor Retail Sales

- We apply a seasonal moving average to the growth rate of the Motor Retail Sales data, with $m = 3$.
- We plot the growth rate series, the estimated seasonality, and the residual after subtracting off the seasonal component.

```
Ret441 <- read.table("retail441.b1",header=TRUE,skip=2)[,2]
Ret441 <- ts(Ret441,start = 1992,frequency=12)
ret.diff <- diff(log(Ret441))


pseas <- 3
period <- 12
seas.smoother <- c(1,rep(0,period-1))
seas.smoother <- c(rep(seas.smoother,pseas),rep(seas.smoother,pseas),1)/(2*pseas+1)
seas.smooth <- stats::filter(ret.diff,seas.smoother,method="convolution")
seas.resid <- ret.diff - seas.smooth
```



```

par(oma=c(2,0,0,0),mar=c(2,4,2,2)+0.1,mfrow=c(3,1),cex.lab=1.2)
plot(ret.diff,ylab="Motor Growth Rate",xlab="",yaxt="n",xaxt="n")
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
plot(ts(seas.smooth,start=1992,frequency=12),xlab="",ylab="Seasonal",yaxt="n",xaxt="n")
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
plot(ts(seas.resid,start=1992,frequency=12),xlab="",ylab="Residual",yaxt="n",xaxt="n")
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
mtext(text="Year",side=1,line=1,outer=TRUE)

```



Remark 3.5.9. Seasonal Adjustment

- Removal of seasonality is called *seasonal adjustment*.
- The *seasonal aggregation* filter is defined as

$$U(B) = 1 + B + B^2 + \dots + B^{s-1}.$$


It annihilates any periodic pattern, and is a basic example of a seasonal adjustment filter.

Example 3.5.10. Seasonal Adjustment of Mauna Loa Growth Rate

- We apply the seasonal aggregation filter to the growth rate of the Mauna Loa CO2 time series.
- This uses a trick: applying $U(B)$ and $1 - B$ together amounts to applying $1 - B^s$, because

$$U(B)(1 - B) = 1 - B^s.$$

```
seas.elim <- diff(mau, lag=12)
plot(seas.elim, xlab="Year", ylab="Mauna Loa")
```



Lesson 3-6: Trend and Seasonality Together

- We may want to extract both trend and seasonality from time series.
- This is more challenging than just extracting trend or seasonality alone.

Proposition 3.6.5

- Stable seasonality has an exactly periodic pattern.
- A trend extraction filter $\Psi(B)$ eliminates stable seasonality if it can be written as

$$\Psi(B) = \Theta(B)U(B),$$

for some filter $\Theta(B)$, where $U(B)$ is the seasonal aggregation filter.

Example 3.6.7. Trend Filters that Eliminate Seasonality

- For s even, consider the trend filter

$$\Psi(B) = \frac{1}{2s}B^{-s/2} + \frac{1}{s}B^{-s/2+1} + \dots + \frac{1}{s}B^{s/2-1} + \frac{1}{2s}B^{s/2}.$$

- By polynomial multiplication, we can show that

$$\Psi(B) = \frac{1}{2s}U(B)(1+B)B^{-s/2}.$$

- Setting $\Theta(B) = (1 + B)B^{-s/2}/(2s)$, we see that $\Psi(B)$ eliminates stable seasonality.

Paradigm 3.6.8. Classical Decomposition from Trend Filtering

- Consider a time series with both trend and seasonal effects.
- We can design a trend filter that eliminates seasonality, obtaining \hat{T}_t .
- Then apply a seasonal filter to $X_t - \hat{T}_t$, the de-trended data, obtaining \hat{S}_t .
- Removing \hat{S}_t from the de-trended data yields the irregular \hat{I}_t .
- Then we have the *classical decomposition*

$$X_t = \hat{T}_t + \hat{S}_t + \hat{I}_t.$$

Example 3.6.13. Classical Decomposition of Western Housing Starts via Linear Filtering

- Consider the trend filter given in Example 3.6.7.
- We construct the classical decomposition, applied to Western Housing Starts

```
compSmooth <- function(data,period,pseas)
{
  # compSmooth
  # Estimates trend, seasonal, and irregular from input data
  # using weighted moving averages.
  # data: an input time series
  # period: observations per year
  # pseas: odd integer giving number of years in seasonal moving average
  #

  n <- length(data)
  trendfilter <- seq(1,period)
  trendfilter <- c(trendfilter[1:(period-1)],rev(trendfilter))/period^2
  seasfilter <- c(1,rep(0,period-1))
  seasfilter <- c(rep(seasfilter,pseas),rep(seasfilter,pseas),1)/(2*pseas+1)
  trend <- stats::filter(data,trendfilter,method="convolution",sides=2)
  trend <- trend[period:(n-period+1)]
  resid <- data[period:(n-period+1)] - trend
  seasonal <- stats::filter(resid,seasfilter,method="convolution",sides=2)
  seasonal <- seasonal[(pseas*period + 1):(length(resid) - pseas*period)]
  irreg <- resid[(pseas*period + 1):(length(resid) - pseas*period)] - seasonal

  return(cbind(trend[(pseas*period + 1):(length(resid) - pseas*period)],seasonal,irreg))
}

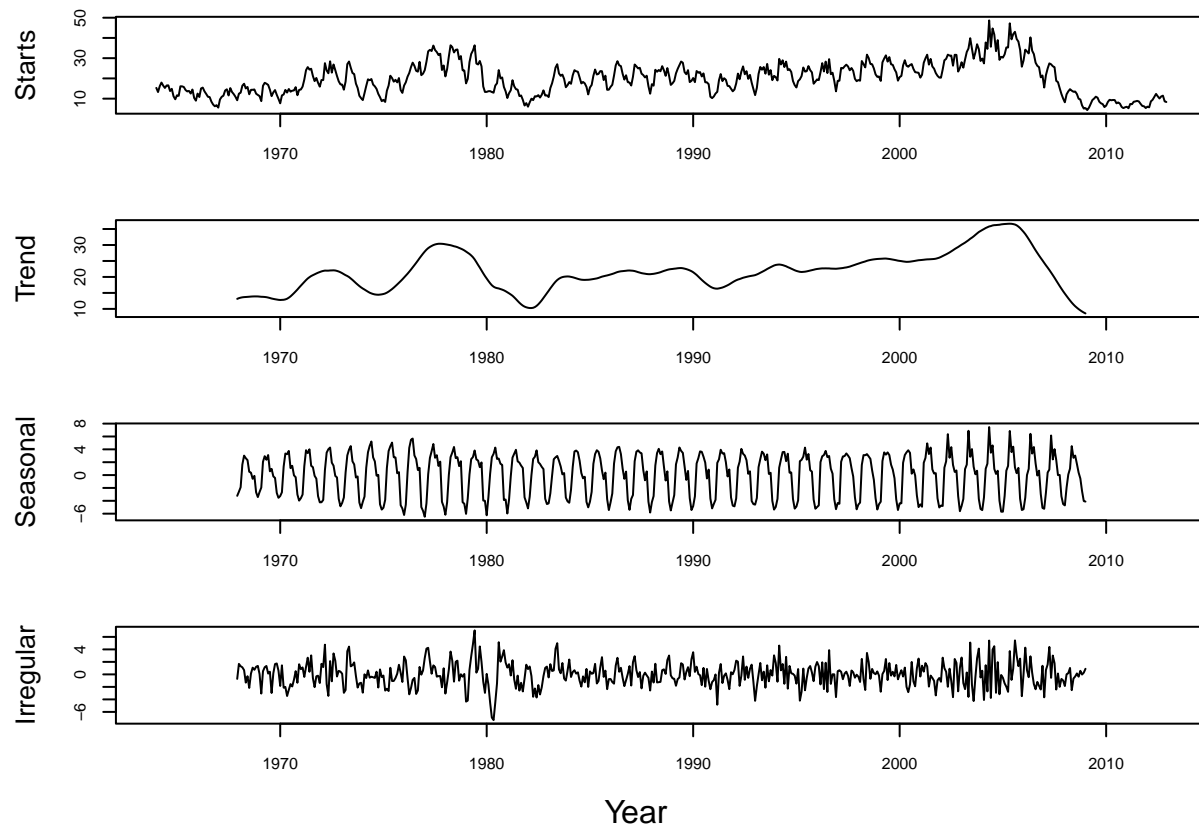
Wstarts <- read.table("Wstarts.b1",skip=2)[,2]
Wstarts <- ts(Wstarts,start = 1964,frequency=12)
pseas <- 3
comps <- compSmooth(Wstarts,12,pseas)

par(oma=c(2,0,0,0),mar=c(2,4,2,2)+0.1,mfrow=c(4,1),cex.lab=1.2)
plot(Wstarts,ylab="Starts",xlab="",yaxt="n",xaxt="n")
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
plot(ts(c(rep(NA,47),comps[,1],rep(NA,47)),start=1964,frequency=12),xlab="",ylab="Trend",yaxt="n",xaxt="n")
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
```

```

plot(ts(c(rep(NA,47),comps[,2],rep(NA,47)),start=1964,frequency=12),xlab="",ylab="Seasonal",yaxt="n",xaxp=c(1,1,1),
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
plot(ts(c(rep(NA,47),comps[,3],rep(NA,47)),start=1964,frequency=12),xlab="",ylab="Irregular",yaxt="n",xaxp=c(1,1,1),
axis(1,cex.axis=.8)
axis(2,cex.axis=.8)
mtext(text="Year",side=1,line=1,outer=TRUE)

```



Lesson 3-7: Integrated Processes

- Sometimes non-stationarity arises in the mean function $\mathbb{E}[X_t] = \mu_t$.
- But the mean may not explain all the non-stationarity of a time series.
- *Integrated Processes* are stochastic processes that exhibit a type of non-stationarity.

Example 3.7.1. Random Walk

- We revisit the *random walk* $\{X_t\}$, where

$$X_t = X_{t-1} + Z_t$$

for $t \geq 1$, and $\{Z_t\}$ are i.i.d. $(0, \sigma^2)$.

- We can initialize with $X_0 = 0$.
- Note that differencing yields the increment:

$$(1 - B)X_t = Z_t.$$

- Writing the recursion out yields

$$X_t = X_0 + \sum_{k=1}^t Z_k.$$

- This is a cumulation, hence an *integrated process*.

Integrated Process

- Suppose that $\{X_t\}$ is a non-stationary stochastic process such that

$$(1 - B)X_t = Z_t$$


is stationary. Then it is called an *integrated process*.

- Differencing $1 - B$ cancels out the cumulation.

Example of Integrated Moving Average

- Suppose that $\{Z_t\}$ is an MA(1), and $X_t = X_{t-1} + Z_t$.
- We simulate a sample path.

```
set.seed(777)
n <- 100
eps <- rnorm(n+1)      # Gaussian input
theta <- .8
z <- eps[-1] + theta*eps[-(n+1)]
x <- cumsum(z)
plot(ts(cbind(x,z)),xlab="Time",ylab="",main="")
```



Example 3.7.3. Unit Root Process

- The AR(1) process $\{X_t\}$ satisfies

$$X_t = \phi X_{t-1} + Z_t,$$

and can be written as

$$(1 - \phi B)X_t = Z_t.$$

So when $\phi = 1$, the AR(1) is a random walk. The polynomial $1 - \phi z$ has root $1/\phi$, and when this root equals one, the root is *unit*. Hence a random walk (and more generally, an integrated process) is sometimes called a *unit root process*.

- Consider a sample path of an AR(1) with high value of ϕ .

```
n <- 100
for(i in 1:2)
{
  set.seed(123)
  z <- rnorm(n)
  x <- rep(0,n)
  phi <- .80+(i-1)*.199
  x0 <- 0
  x[1] <- x0 + z[1]
  for(t in 2:n) { x[t] <- phi*x[t-1] + z[t] }
  plot(ts(x),xlab="Time",ylab="")
}
```

