

Time Series: A First Course with Bootstrap Starter

Contents

Lesson 6: Active Recall space	3
Lesson 6: Takes and tools	4
Takes	4
Questions	4
Links to JD+	4
Overview in book	4
Concepts	4
R skills	4
Lesson 6-1: Spectral Density	4
Definition 6.1.2.	4
Remark 6.1.6. Spectral Representation of the Autocovariance.	4
Exercise 6.4. MA(1) Spectral Density.	4
Fact 6.1.8. Further Properties of the Spectral Density	5
Exercise 6.7. MA(q) Spectral Density Computation.	5
Corollary 6.1.9.	6
Theorem 6.1.12.	7
Exercise 6.12. ARMA(p,q) Spectral Density.	7
Corollary 6.1.14.	8
Theorem 6.1.16. MA(∞) Representation.	8
Corollary 6.1.17. AR(∞) Representation.	8
Lesson 6-2: Filtering in Frequency Domain	8
Example 6.2.1. Business Cycle in Housing Starts.	9
Remark 6.2.2. Spectral Peaks and Oscillation Frequencies	12
Fact 6.2.5. Suppression and Extraction	12
Exercise 6.22. The Ideal Low-Pass Filter	12
Exercise 6.23. The Ideal Band-Pass Filter	13
Example 6.2.7. The Hodrick-Prescott Filter	14
Lesson 6-3: Inverse Autocovariance	16
Paradigm 6.3.1. Whitening a Time Series	16
Definition 6.3.2	16
Example 6.3.3. Prediction of an MA(1) from an Infinite Past	16
Definition 6.3.4	17
Example 6.3.6. The Inverse Autocovariance of an MA(1)	17
Exercise 6.34. Inverse Autocovariances of an AR(1)	18
Example 6.3.8. Inverse ACF and Optimal Interpolation	18
Fact 6.4.2. Spectral Representation of a Symmetric Matrix	18
Definition 6.4.3. Fourier Frequencies	18
Theorem 6.4.5. Spectral Decomposition of Toeplitz Covariance Matrices	18
Remark 6.4.11. Positive Spectral Density	19
Exercise 6.43. Eigenvalues of an MA(1) Toeplitz Matrix.	19
Exercise 6.45. Eigenvalues of an MA(1) Inverse Toeplitz Matrix.	20

Lesson 6-5: Partial Autocorrelation	21
Definition 6.5.1.	21
Example 6.5.2. Partial Autocorrelation of an $AR(p)$ Process	21
Proposition 6.5.5.	21
Exercise 6.54. PACF of $MA(q)$	21
Exercise 6.55. PACF pf $ARMA(p,q)$	23
Lesson 6-6: AR and MA Identification	24
Paradigm 6.6.1. Characterizing AR and MA Processes	24
ACF	24
IACF	24
PACF	25
Finding Truncation	25
Example 6.6.2. $MA(3)$ Identification	25
Example 6.6.3. $AR(4)$ Identification	29
Paradigm 6.6.7. Identification by Whitening	32
Example 6.6.8. $AR(p)$ Whitening Models	32
Exercise 6.61. Whitening an $AR(p)$ Process	32
Lesson 7-2: Discrete Fourier Transform	43
Definition 7.2.3.	43
Definition 7.2.4.	44
Proposition 7.2.7.	44
Exercise 7.14. DFT of an $AR(1)$	44
Corollary 7.2.9. Decorrelation Property of the DFT.	47
Exercise 7.18. DFT of Wolfer Sunspots.	47
Exercise 7.20. DFT of Mauna Loa Growth Rate.	50
Lesson 7-3: Spectral Representation	53
Definition 7.3.1.	53
Paradigm 7.3.4. Time Series Defined as a Stochastic Integral	53
Corollary 7.3.8.	54
Example 7.3.10. Time Shift	54
Definition 7.3.11.	54
Fact 7.3.12. Action of Phase Delay.	54
Example 7.3.13. Simple Moving Average Filters Cause Delay	55
Example 7.3.14. Differencing Causes an Advance	55
Exercise 7.29. Phase and Gain for Simple Moving Average.	55
Exercise 7.30. Phase and Gain for Seasonal Aggregation Filter.	60

Lesson 6: Active Recall space

Lesson 6: Takes and tools

Takes

Questions

key = small TPs on my data

Links to JD+

Overview in book

Concepts

R skills

Lesson 6-1: Spectral Density

- We define the spectral density, which allows us to do time series analysis in the frequency (or Fourier) domain.

Definition 6.1.2.

- The *spectral density* of a stationary time series is

$$f(\lambda) = \sum_{k=-\infty}^{\infty} \gamma(k) e^{-i\lambda k},$$

for $\lambda \in [-\pi, \pi]$.

- Also the spectral density is the restriction of the AGF to the unit circle: $f(\lambda) = G(e^{-i\lambda})$.
- A sufficient condition for existence is absolute summability of the autocovariances. This also guarantees $f(\lambda)$ is continuous.

Remark 6.1.6. Spectral Representation of the Autocovariance.

- By Fourier inversion, we can recover the autocovariances from the spectral density:

$$\gamma(k) = \frac{1}{2\pi} \int_{-\pi}^{\pi} f(\lambda) e^{i\lambda k} d\lambda.$$

- So for $k = 0$, we see the process' variance is the average integral of f .


Exercise 6.4. MA(1) Spectral Density.

- For an MA(1), the spectral density is

$$f(\lambda) = \gamma(0) + 2\gamma(1) \cos(\lambda) = \gamma(0) (1 + 2\rho(1) \cos(\lambda)).$$

- We plot $f(\lambda)$ for $\rho(1) = .4$ and $\gamma(0) = 1$.
- The units are in π , called “Cycles”.

```
mesh <- 1000
lambda <- pi*seq(-mesh,mesh)/mesh
rho.1 <- .4
spec <- 1 + 2*rho.1*cos(lambda)
plot(ts(spec,start=-1,frequency=mesh),xlab="Cycles",ylab="")
abline(h=0,col=2)
```



Fact 6.1.8. Further Properties of the Spectral Density

- Because $\gamma(k) = \gamma(-k)$, $f(\lambda)$ is real and even.
- Also $f(\lambda) \geq 0$ follows from positive definite property.

Exercise 6.7. MA(q) Spectral Density Computation.

- For an MA(q), the spectral density is

$$f(\lambda) = \gamma(0) \left(1 + 2 \sum_{k=1}^q \rho(k) \cos(\lambda k) \right).$$

- Since it is even, we can just focus on $\lambda \in [0, \pi]$.

```
maq.spec <- function(ma.acf, mesh)
{
  q <- length(ma.acf)-1
  lambda <- pi*seq(0, mesh)/mesh
  spec <- ma.acf[1]*cos(0*lambda)
  if(q > 0)
  {
    for(k in 1:q)
    {
      spec <- spec + 2*ma.acf[k+1]*cos(k*lambda)
    }
  }
}
```

```

    return(spec)
}


```

- We plot $f(\lambda)$ with $q = 2$, $\gamma(0) = 1$, $\rho(1) = 0$, and $\rho(2) = .6$.
- However, these values do not correspond to a positive definite autocovariance, and the resulting function takes negative values.
- Q: est ce grave ? see recording
- arguments fonction: detail ?
- variance et q autocorrelations (1 à q)
- ma.acf stocke ces q+1 valeurs, q est deduit

```

spec <- maq.spec(c(1,0,.6),mesh)
plot(ts(spec,start=0,frequency=mesh),xlab="Cycles",ylab="")
abline(h=0,col=2)

```



Corollary 6.1.9.

- Suppose we filter stationary $\{X_t\}$ with some $\psi(B)$, yielding $Y_t = \psi(B)X_t$. Then the spectral densities of input and output are related by

$$f_y(\lambda) = |\psi(e^{-i\lambda})|^2 f_x(\lambda).$$

- This follows from Theorem 5.6.6.
- We call $\psi(e^{-i\lambda})$ the *frequency response function* of the filter $\psi(B)$.
- We call $|\psi(e^{-i\lambda})|^2$ the *squared gain function* of the filter $\psi(B)$.

Theorem 6.1.12.

Let $\{X_t\}$ be a stationary ARMA(p, q) process such that $\phi(B)X_t = \theta(B)Z_t$, for $Z_t \sim \text{WN}(0, \sigma^2)$. Suppose $\phi(z)$ has no roots on the unit circle. Then the spectral density exists:

$$f(\lambda) = \sigma^2 \frac{|\theta(e^{-i\lambda})|^2}{|\phi(e^{-i\lambda})|^2}$$

Exercise 6.12. ARMA(p, q) Spectral Density.

- We write code for the ARMA spectral density, based on the formula of Theorem 6.1.12, taking as input the $\theta(z)$ and $\phi(z)$ polynomials.


```
armapq.spec <- function(ar.coef,ma.coef,sigma,mesh)
{
  p <- length(ar.coef)
  q <- length(ma.coef)
  lambda <- pi*seq(0,mesh)/mesh
  spec.ar <- rep(1,mesh+1)
  if(p > 0)
  {
    for(k in 1:p)
    {
      spec.ar <- spec.ar - ar.coef[k]*exp(-1i*lambda*k)
    }
  }
  spec.ma <- rep(1,mesh+1)
  if(q > 0)
  {
    for(k in 1:q)
    {
      spec.ma <- spec.ma + ma.coef[k]*exp(-1i*lambda*k)
    }
  }
  spec <- sigma^2*Mod(spec.ma)^2/Mod(spec.ar)^2
  return(spec)
}
```

- We plot the spectral density of the cyclic ARMA(2,1) process of Example 5.7.2: for $\rho \in (0, 1)$ and $\omega \in (0, \pi)$, let $\{X_t\}$ satisfy

$$(1 - 2\rho \cos(\omega)B + \rho^2 B^2)X_t = (1 - \rho \cos(\omega)B)Z_t.$$

- We set $\rho = .8$ and $\omega = \pi/6$.

```
spec <- NULL
mesh <- 1000
rho <- .8
omega <- pi/6
ar.coef <- c(2*rho*cos(omega), -1*rho^2)
ma.coef <- -1*rho*cos(omega)
spec <- armapq.spec(ar.coef,ma.coef,1,mesh)
plot(ts(spec,start=0,frequency=mesh),xlab="Cycles",ylab="",main="")
```



Corollary 6.1.14.

- Let $\{X_t\}$ be a weakly stationary, mean zero time series with strictly positive spectral density of form given in Theorem 6.1.12. Then there exists a white noise $\{Z_t\}$ such that $\phi(B)X_t = \theta(B)Z_t$.
- This is proved by defining $Z_t = \psi(B)X_t$ with $\psi(z) = \phi(z)/\theta(z)$, and checking that $\{Z_t\}$ is white noise.
- This $\psi(B)$ is a *whitening filter*. It transforms a time series to white noise!

Theorem 6.1.16. MA(∞) Representation.

Let $\{X_t\}$ be a weakly stationary, mean zero time series with autocovariance function $\gamma(k)$ that is absolutely summable, and positive spectral density. Then $\{X_t\}$ is an MA(∞) process with respect to some white noise $\{Z_t\}$:

$$X_t = \sum_{j \geq 0} \psi_j Z_{t-j},$$

and $\psi_0 = 1$.

Corollary 6.1.17. AR(∞) Representation.

Under the assumptions of Theorem 6.1.16, $\{X_t\}$ is an AR(∞) process with respect to the same white noise $\{Z_t\}$:

$$X_t = - \sum_{j \geq 1} \pi_j X_{t-j} + Z_t.$$

Lesson 6-2: Filtering in Frequency Domain

- Filters extract (or suppress) features of interest from a time series.

- Using Corollary 6.1.9, we can see in frequency domain how extraction and suppression occurs.

Example 6.2.1. Business Cycle in Housing Starts.

- Example 3.6.13 decomposed West Housing Starts into Trend, Seasonal, and Irregular.

```
hpsa <- function(n,period,q,r)
{
  # hpsa
  #   gives an HP filter for seasonal data
  #   presumes trend+seas+irreg structure
  #   trend is integrated rw
  #   seas is seasonal rw
  #   irreg is un
  #   q is snr for trend to irreg
  #   r is snr for seas to irreg

  # define trend differencing matrix

  delta.mat <- diag(n)
  temp.mat <- 0*diag(n)
  temp.mat[-1,-n] <- -2*diag(n-1)
  delta.mat <- delta.mat + temp.mat
  temp.mat <- 0*diag(n)
  temp.mat[c(-1,-2),c(-n,-n+1)] <- 1*diag(n-2)
  delta.mat <- delta.mat + temp.mat
  diff.mat <- delta.mat[3:n,]

  # define seasonal differencing matrix

  delta.mat <- diag(n)
  temp.mat <- 0*diag(n)
  inds <- 0
  for(t in 1:(period-1))
  {
    temp.mat <- 0*diag(n)
    temp.mat[-(1+inds),-(n-inds)] <- 1*diag(n-t)
    delta.mat <- delta.mat + temp.mat
    inds <- c(inds,t)
  }
  sum.mat <- delta.mat[period:n,]

  # define two-comp sig ex matrices

  #trend.mat <- solve(diag(n) + t(diff.mat) %*% diff.mat/q)
  #seas.mat <- solve(diag(n) + t(sum.mat) %*% sum.mat/r)
  trend.mat <- diag(n) - t(diff.mat) %*% solve(q*diag(n-2) + diff.mat %*%
    t(diff.mat)) %*% diff.mat
  seas.mat <- diag(n) - t(sum.mat) %*% solve(r*diag(n-period+1) + sum.mat %*%
    t(sum.mat)) %*% sum.mat

  # define three-comp sig ex matrices

  trend.filter <- solve(diag(n) - trend.mat %*% seas.mat) %*%
```

```

trend.mat %*% (diag(n) - seas.mat)
seas.filter <- solve(diag(n) - seas.mat %*% trend.mat) %*%
  seas.mat %*% (diag(n) - trend.mat)
irreg.filter <- diag(n) - (trend.filter + seas.filter)


filters <- list(trend.filter,seas.filter,irreg.filter)
return(filters)
}

Wstarts <- read.table("Wstarts.b1",skip=2)[,2]
Wstarts <- ts(Wstarts,start = 1964,frequency=12)
n <- length(Wstarts)
q <- .0001
r <- 1
hp.filters <- hpsa(n,12,q,r)

wstarts.trend <- ts(hp.filters[[1]] %*% Wstarts,start=1964,frequency=12)
wstarts.seas <- ts(hp.filters[[2]] %*% Wstarts,start=1964,frequency=12)
wstarts.irreg <- ts(hp.filters[[3]] %*% Wstarts,start=1964,frequency=12)


par(oma=c(2,0,0,0),mar=c(2,4,2,2)+0.1,mfrow=c(4,1),cex.lab=.8)
plot(Wstarts, ylab="West Starts",xlab="",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
plot(wstarts.trend,xlab="",ylab = "Trend",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
plot(wstarts.seas,xlab="",ylab = "Seasonal",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
plot(wstarts.irreg,xlab="",ylab = "Irregular",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
mtext(text="Year",side=1,line=1,outer=TRUE)

```



- We fit (e.g. via ordinary least squares) an AR(26) to the Irregular component, and plot the log spectral density.
- The units are in terms of $2\pi/12$, so the x-axis numbers represent multiples of $\pi/6$.
- There are vertical bands for frequency between .5 and .1, corresponding to period between 2 and 10 years. This is the *business cycle* range.

```
ar.fit <- spec.ar(ts(wstarts.irreg,frequency=12),plot=FALSE)
plot(ts(log(ar.fit$spec),start=0,frequency=500/6),xlab="Frequency per Year",
      ylab="Log Spectrum",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
abline(v=.5,lty=3)
abline(v=.1,lty=3)
```



Remark 6.2.2. Spectral Peaks and Oscillation Frequencies

- Higher values of the spectral density correspond to frequencies with more variability.
- A peak in the spectral density at a frequency λ corresponds to an oscillation, or cyclical effect in the process.

Fact 6.2.5. Suppression and Extraction

- Suppose $\{X_t\}$ is filtered with $\psi(B)$, so that $Y_t = \psi(B)X_t$.
- If $\psi(e^{-i\lambda}) = 0$, then $f_y(\lambda) = 0$, and λ is *suppressed*. The set of such frequencies is the *stop-band*.
- If $\psi(e^{-i\lambda}) = 1$, then $f_y(\lambda) = f_x(\lambda)$, and λ is *extracted*. The set of such frequencies is the *pass-band*.


Exercise 6.22. The Ideal Low-Pass Filter

- The Ideal Low-pass is defined with frequency response function

$$\psi(e^{-i\lambda}) = \begin{cases} 1 & \text{if } |\lambda| \leq \mu \\ 0 & \text{else.} \end{cases}$$

- We plot with cut-off $\mu = \pi/5$.

```
mu <- pi/5
mesh <- 1000
lambda <- pi*seq(0,mesh)/mesh
psi.frf <- rep(0,mesh+1)
psi.frf[lambda <= mu] <- 1
plot(ts(psi.frf,start=0,frequency=mesh),ylab="",xlab="Cycles")
```



- It is hard to implement, since the filter coefficients decay slowly (and hence truncation is expensive).


Exercise 6.23. The Ideal Band-Pass Filter

- The Ideal Band-pass is defined with frequency response function

$$\psi(e^{-i\lambda}) = \begin{cases} 1 & \text{if } \mu_1 < |\lambda| \leq \mu_2 \\ 0 & \text{else.} \end{cases}$$

- We plot with cut-offs $\mu_1 = \pi/60$ and $\mu_2 = \pi/12$.

```
mu1 <- pi/60
mu2 <- pi/12
mesh <- 1000
lambda <- pi*seq(0,mesh)/mesh
psi.frf <- rep(0,mesh+1)
psi.frf[lambda <= mu2] <- 1
psi.frf[lambda <= mu1] <- 0
plot(ts(psi.frf,start=0,frequency=mesh),ylab="",xlab="Cycles")
```



Example 6.2.7. The Hodrick-Prescott Filter

- Proposed by Whitaker in 1927, but known as Hodrick-Prescott, the filter does trend extraction.
- The frequency response function resembles an ideal low-pass.
- For the decomposition of West Housing Starts, we used a modified Hodrick-Prescott that is adapted for finite samples, and accounts for seasonality.
- The Hodrick-Prescott filter depends on a parameter $q > 0$:

$$\psi(e^{-i\lambda}) = \frac{q}{q + |1 - e^{-i\lambda}|^4}.$$


- There is a formula for the coefficients. We plot the filter coefficients with $q = 1/1600$.

```
q <- 1/1600

s <- (2*q + 2*q^(1/2)*(q+16)^(1/2))^(1/2)
r <- (q^(1/2) + (q+16)^(1/2) + s)/4
c <- q/r^2
phi1 <- 2*(q^(1/2)-(q+16)^(1/2))/(4*r)
phi2 <- (q^(1/2)+(q+16)^(1/2) - s)/(4*r)
theta <- atan(s/4)

lags <- seq(0,100)
psi <- 2*c*r^(4-lags)*sin(theta)*(r^2*sin(theta*(1+lags)) - sin(theta*(lags-1)))
psi <- psi/((1-2*r^2*cos(2*theta)+r^4)*(r^2-1)*(1-cos(2*theta)))
psi <- c(rev(psi),psi[-1])
```


```
plot(ts(psi,start=-101),xlab="Index",ylab="",yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
```



- Next, we show the squared gain function.

```
grid <- 1000
lambda <- pi*seq(0,grid+1)/grid
gain <- q/(q + (2 - 2*cos(lambda))^2)
sq.gain <- gain^2

plot(ts(sq.gain,start=0,frequency=grid),xlab="Frequency",ylab="Squared Gain",
     yaxt="n",xaxt="n")
axis(1,cex.axis=.5)
axis(2,cex.axis=.5)
```



Lesson 6-3: Inverse Autocovariance

- Inverse autocovariances are related to whitening filters, and can be used for model identification.

Paradigm 6.3.1. Whitening a Time Series

- A *whitening filter* reduces a time series to white noise.
- Suppose $\{X_t\}$ is stationary with positive spectral density $f(\lambda)$. Then a whitening filter $\psi(B)$ has squared gain function

$$|\psi(e^{-i\lambda})|^2 \propto 1/f(\lambda).$$

- So $\psi(B)$ depends on the spectral density of the time series we are whitening. We can find $\psi(B)$ causal, i.e., $\psi(z)$ is a power series.

Definition 6.3.2

- A weakly stationary process is *invertible* if its spectral density is positive.
- By Corollary 6.1.17, the process has an $\text{AR}(\infty)$ representation, so we can “invert” the time series into a white noise.
- For prediction problems, a process should be invertible.

Example 6.3.3. Prediction of an MA(1) from an Infinite Past

- Let $\{X_t\}$ be an *invertible* MA(1) process with MA polynomial $1 + \theta_1 z$.
- Suppose we want to forecast 1-step ahead: we seek $\hat{X}_{t+1} = P_{\overline{\text{sp}}\{X_s, s \leq t\}}[X_{t+1}]$.
- This forecast is a causal filter: $\hat{X}_{t+1} = \sum_{j \geq 0} \psi_j X_{t-j}$, with ψ_j to be determined from normal equations.

- The normal equations give us, for any $h \geq 0$:

$$\gamma(h+1) = \text{Cov}[X_{t+1}, X_{t-h}] = \text{Cov}[\hat{X}_{t+1}, X_{t-h}] = \sum_{j \geq 0} \psi_j \text{Cov}[X_{t-j}, X_{t-h}] = \sum_{j \geq 0} \psi_j \gamma(h-j).$$

- To solve this, rewrite the right hand side using Fourier inversion:

$$\sum_{j \geq 0} \psi_j \gamma(h-j) = \sum_{j \geq 0} \psi_j (2\pi)^{-1} \int_{-\pi}^{\pi} e^{i\lambda(h-j)} f(\lambda) d\lambda = (2\pi)^{-1} \int_{-\pi}^{\pi} e^{i\lambda h} \sum_{j \geq 0} \psi_j e^{-i\lambda j} f(\lambda) d\lambda = (2\pi)^{-1} \int_{-\pi}^{\pi} e^{i\lambda h} \psi(e^{-i\lambda}) f(\lambda) d\lambda.$$

- Recall that $f(\lambda) = \sigma^2 |1 + \theta_1 e^{-i\lambda}|^2$. The invertibility assumption means that $1 + \theta_1 e^{-i\lambda}$ is non-zero for all λ .
- Claim: the prediction filter has frequency response function

$$\psi(e^{-i\lambda}) = \frac{\theta_1}{1 + \theta_1 e^{-i\lambda}},$$

which is well-defined by the invertibility assumption.

- To prove this claim, we plug in and check! The right hand side becomes

$$\sum_{j \geq 0} \psi_j \gamma(h-j) = (2\pi)^{-1} \int_{-\pi}^{\pi} e^{i\lambda h} \frac{\theta_1}{1 + \theta_1 e^{-i\lambda}} f(\lambda) d\lambda = (2\pi)^{-1} \int_{-\pi}^{\pi} e^{i\lambda h} \theta_1 (1 + \theta_1 e^{i\lambda}) \sigma^2 d\lambda = 1_{\{h=0\}} \theta_1 \sigma^2.$$

This is the same as $\gamma(h+1)$ for $h \geq 0$, so the claim is true!

- To get the coefficients:

$$\psi(z) = \theta_1 (1 + \theta_1 z)^{-1} = \theta_1 \sum_{j \geq 0} (-\theta_1)^j z^j.$$

So $\psi_j = \theta_1^{j+1} (-1)^j$.

Definition 6.3.4

- Suppose $\{X_t\}$ is an invertible weakly stationary time series with autocovariance $\gamma(h)$. Then the *inverse autocovariance* is the sequence $\xi(k)$ such that

$$\sum_{k=-\infty}^{\infty} \gamma(k) \xi(j-k) = 1_{\{j=0\}}.$$

- The inverse autocorrelation is $\zeta(k) = \xi(k)/\xi(0)$.
- We can compute the inverse autocovariance from the spectral density:

$$\xi(k) = \frac{1}{2\pi} \int_{-\pi}^{\pi} e^{i\lambda k} \frac{1}{f(\lambda)} d\lambda.$$

Example 6.3.6. The Inverse Autocovariance of an MA(1)

- Consider an MA(1) process with $\theta_1 \in (-1, 1)$, which implies it is invertible.
- So the inverse autocovariance is

$$\xi(k) = \frac{1}{2\pi} \int_{-\pi}^{\pi} e^{i\lambda k} \sigma^{-2} |1 + \theta_1 e^{-i\lambda}|^{-2} d\lambda.$$

- This resembles the autocovariance of an AR(1), with parameter $\phi_1 = -\theta_1$, and input variance σ^{-2} .
- So by using the formula for AR(1) autocovariance, we find

$$\xi(k) = \sigma^{-2} \frac{(-\theta_1)^{|k|}}{1 - \theta_1^2}.$$

Exercise 6.34. Inverse Autocovariances of an AR(1)

- Consider the AR(1) process with $\phi(z) = 1 - \phi_1 z$. What are the inverse autocovariances?
- So the inverse autocovariance is

$$\xi(k) = \frac{1}{2\pi} \int_{-\pi}^{\pi} e^{i\lambda k} \sigma^2 |1 - \phi_1 e^{-i\lambda}|^2 d\lambda.$$

- This resembles the autocovariance of an MA(1), with parameter $\theta_1 = -\phi_1$, and input variance σ^2 .
- So by using the formula for MA(1) autocovariance, we find

$$\xi(k) = \sigma^2 \begin{cases} 1 + \phi_1^2 & \text{if } k = 0 \\ -\phi_1 & \text{if } k = \pm 1 \\ 0 & \text{if } |k| > 1. \end{cases}$$

Example 6.3.8. Inverse ACF and Optimal Interpolation

- Suppose $\{X_t\}$ is stationary, mean zero, and invertible.
- Suppose that X_0 is missing. What is the optimal estimator?
- We seek $\widehat{X}_0 = P_{\overline{\text{SP}}\{X_j, j \neq 0\}}[X_0]$, which is a linear filter $\psi(B)$ of the data, such that $\psi_0 = 0$.
- Claim: $\psi_j = -\zeta(j)$ for $j \neq 0$, and $\psi_0 = 0$.
- Proof: check the normal equations. First $X_0 - \widehat{X}_0 = \sum_j \zeta(j) X_{-j}$, since $\zeta(0) = 1$. The covariance of this with any X_{-k} for $k \neq 0$ is

$$\text{Cov}[X_0 - \widehat{X}_0, X_{-k}] = \sum_j \zeta(j) \gamma(k - j),$$

which is zero (since $k \neq 0$) by definition of inverse autocovariance. This verifies the normal equations!

Lesson 6-4: Toeplitz Matrices

- We discuss a decomposition of Toeplitz matrices, with connections to the spectral density.
- This is useful for model fitting and prediction.

Fact 6.4.2. Spectral Representation of a Symmetric Matrix

- Let A^* denote the conjugate transpose of a matrix A .
- A matrix A is Hermitian if $A^* = A$.
- A matrix U is unitary if $U^{-1} = U^*$.
- For any Hermitian A , there exists unitary U such that $A = UDU^*$, where D is diagonal with real entries.
- The columns of U are the eigenvectors of A , and D has the eigenvalues of A .

Definition 6.4.3. Fourier Frequencies

- For any n , the *Fourier frequencies* are defined as $\lambda_\ell = 2\pi\ell/n$, for $[n/2] - n + 1 \leq \ell \leq [n/2]$.
- When n is even, this excludes the frequencies π and $-\pi$ both being in the set, since these are redundant.
- We define an $n \times n$ -dimensional matrix Q , whose entries are complex exponentials evaluated at Fourier frequencies:

$$Q_{jk} = n^{-1/2} e^{ij\lambda_{[n/2]-n+k}}.$$

- The matrix Q is unitary, and can be used in an approximation result for Toeplitz matrices.

Theorem 6.4.5. Spectral Decomposition of Toeplitz Covariance Matrices

- Let Γ_n denote the autocovariance matrix of a sample of size n from a stationary time series. Suppose the autocovariances $\gamma(k)$ are absolutely summable. Then

$$\Gamma_n \approx Q\Lambda Q^*,$$

with Λ diagonal with entries $f(\lambda_{[n/2]-n+k})$. The approximation \approx means that the difference entry-by-entry tends to zero as n tends to ∞ .

- It can also be shown that $\Lambda \approx Q^* \Gamma_n Q$.
- If the process is invertible, then $\Gamma_n^{-1} \approx Q \Lambda^{-1} Q^*$ as well.

Remark 6.4.11. Positive Spectral Density

Since the eigenvalues of a symmetric non-negative definite matrix Γ_n are real and non-negative, we can show that the spectral density of a stationary process must be non-negative.

Exercise 6.43. Eigenvalues of an MA(1) Toeplitz Matrix.

- Consider an MA(1) with parameter $\theta = .8$.
- Compute the eigenvalues of Γ_n for various n , and compare to the spectral density evaluated at the Fourier frequencies.

```
theta <- .8

n <- 10
lambda <- 2*pi*seq(0,n-1)/n
Gamma <- toeplitz(c(1+theta^2, theta, rep(0,n-2)))
eigen(Gamma)$values

## [1] 3.1751888 2.9860057 2.6877772 2.3046640 1.8677037 1.4122963 0.9753360
## [8] 0.5922228 0.2939943 0.1048112

rev(sort(1+theta^2 + 2*theta*cos(lambda)))

## [1] 3.2400000 2.9344272 2.9344272 2.1344272 2.1344272 1.1455728 1.1455728
## [8] 0.3455728 0.3455728 0.0400000

n <- 20
lambda <- 2*pi*seq(0,n-1)/n
Gamma <- toeplitz(c(1+theta^2, theta, rep(0,n-2)))
eigen(Gamma)$values

## [1] 3.22212932 3.16891649 3.08155019 2.96198204 2.81288299 2.63758368
## [7] 2.44000000 2.22454564 1.99603349 1.75956815 1.52043185 1.28396651
## [13] 1.05545436 0.84000000 0.64241632 0.46711701 0.31801796 0.19844981
## [19] 0.11108351 0.05787068

rev(sort(1+theta^2 + 2*theta*cos(lambda)))

## [1] 3.2400000 3.1616904 3.1616904 2.9344272 2.9344272 2.5804564 2.5804564
## [8] 2.1344272 2.1344272 1.6400000 1.6400000 1.1455728 1.1455728 0.6995436
## [15] 0.6995436 0.3455728 0.3455728 0.1183096 0.1183096 0.0400000

n <- 30
lambda <- 2*pi*seq(0,n-1)/n
Gamma <- toeplitz(c(1+theta^2, theta, rep(0,n-2)))
eigen(Gamma)$values

## [1] 3.23179092 3.20724791 3.16662281 3.11033250 3.03895459 2.95322151
## [7] 2.85401300 2.74234707 2.61936957 2.48634242 2.34463064 2.19568840
## [13] 2.04104405 1.88228444 1.72103867 1.55896133 1.39771556 1.23895595
## [19] 1.08431160 0.93536936 0.79365758 0.66063043 0.53765293 0.42598700
## [25] 0.32677849 0.24104541 0.16966750 0.11337719 0.07275209 0.04820908
```

```
rev(sort(1+theta^2 + 2*theta*cos(lambda)))
```

```
## [1] 3.24000000 3.20503616 3.20503616 3.10167273 3.10167273 2.93442719
## [7] 2.93442719 2.71060897 2.71060897 2.44000000 2.44000000 2.13442719
## [13] 2.13442719 1.80724554 1.80724554 1.47275446 1.47275446 1.14557281
## [19] 1.14557281 0.84000000 0.84000000 0.56939103 0.56939103 0.34557281
## [25] 0.34557281 0.17832727 0.17832727 0.07496384 0.07496384 0.04000000
```

Exercise 6.45. Eigenvalues of an MA(1) Inverse Toeplitz Matrix.

- Consider an MA(1) with parameter $\theta = .8$.
- Compute the eigenvalues of Γ_n^{-1} for various n , and compare to the reciprocal spectral density evaluated at the Fourier frequencies.

```
theta <- .8
```

```
n <- 10
lambda <- 2*pi*seq(0,n-1)/n
Gamma <- toeplitz(c(1+theta^2, theta, rep(0,n-2)))
eigen(solve(Gamma))$values
```

```
## [1] 9.5409612 3.4014259 1.6885536 1.0252877 0.7080667 0.5354168 0.4339027
## [8] 0.3720547 0.3348955 0.3149419
```

```
1/sort(1+theta^2 + 2*theta*cos(lambda))
```

```
## [1] 25.0000000 2.8937462 2.8937462 0.8729257 0.8729257 0.4685098
## [7] 0.4685098 0.3407820 0.3407820 0.3086420
```

```
n <- 20
lambda <- 2*pi*seq(0,n-1)/n
Gamma <- toeplitz(c(1+theta^2, theta, rep(0,n-2)))
eigen(solve(Gamma))$values
```

```
## [1] 17.2799081 9.0022362 5.0390574 3.1444765 2.1407913 1.5566230
## [7] 1.1904762 0.9474593 0.7788365 0.6577079 0.5683213 0.5009936
## [13] 0.4495300 0.4098361 0.3791349 0.3555071 0.3376118 0.3245120
## [19] 0.3155653 0.3103538
```

```
1/sort(1+theta^2 + 2*theta*cos(lambda))
```

```
## [1] 25.0000000 8.4524013 8.4524013 2.8937462 2.8937462 1.4295035
## [7] 1.4295035 0.8729257 0.8729257 0.6097561 0.6097561 0.4685098
## [13] 0.4685098 0.3875283 0.3875283 0.3407820 0.3407820 0.3162865
## [19] 0.3162865 0.3086420
```

```
n <- 30
lambda <- 2*pi*seq(0,n-1)/n
Gamma <- toeplitz(c(1+theta^2, theta, rep(0,n-2)))
eigen(solve(Gamma))$values
```

```
## [1] 20.7429793 13.7453088 8.8201163 5.8938806 4.1485958 3.0601769
## [7] 2.3474895 1.8599359 1.5137056 1.2599892 1.0690964 0.9222441
## [13] 0.8071312 0.7154532 0.6414527 0.5810445 0.5312693 0.4899453
## [19] 0.4554380 0.4265064 0.4021972 0.3817713 0.3646511 0.3503838
## [25] 0.3386133 0.3290605 0.3215090 0.3157938 0.3117938 0.3094260
```

```
1/sort(1+theta^2 + 2*theta*cos(lambda))
```

```
## [1] 25.0000000 13.3397651 13.3397651 5.6076674 5.6076674 2.8937462
## [7] 2.8937462 1.7562623 1.7562623 1.1904762 1.1904762 0.8729257
## [13] 0.8729257 0.6789998 0.6789998 0.5533282 0.5533282 0.4685098
## [19] 0.4685098 0.4098361 0.4098361 0.3689208 0.3689208 0.3407820
## [25] 0.3407820 0.3224067 0.3224067 0.3120090 0.3120090 0.3086420
```

Lesson 6-5: Partial Autocorrelation

- Recall from linear models that partial correlation allows us to explore the relationship between a dependent variable and a covariate, while accounting for other covariates.
- We apply this concept to stationary time series, where we look at the relationship between time present and time past, while accounting for the in-between times.

Definition 6.5.1.

- The partial correlation function (PACF) of stationary time series $\{X_t\}$ is a sequence $\kappa(h)$ defined by $\kappa(1) = \text{Corr}[X_1, X_0]$ and

$$\kappa(h) = \text{Corr}[X_h, X_0 | X_1, \dots, X_{h-1}]$$

when $h \geq 2$. The conditioning stands for projection (of the demeaned time series) on the random variables.

- What does this mean? Linearly predict X_h from X_1, \dots, X_{h-1} , and call that \hat{X}_h . Also linearly predict X_0 from X_1, \dots, X_{h-1} , and call that \hat{X}_0 . Then $\kappa(h)$ is the correlation of the prediction errors:

$$\kappa(h) = \text{Corr}[X_h - \hat{X}_h, X_0 - \hat{X}_0].$$

- Because of stationarity, we could also write

$$\kappa(h) = \text{Corr}[X_{t+h}, X_t | X_{t+1}, \dots, X_{t+h-1}]$$

for any t .

Example 6.5.2. Partial Autocorrelation of an AR(p) Process

- Suppose $\{X_t\}$ is an AR(1), where $\phi(z) = 1 - \phi_1 z$. Then $\kappa(1) = \phi_1$.
- Also for $h \geq 2$, $\hat{X}_h = \phi_1 X_{h-1}$ and $\hat{X}_0 = \phi_1 X_1$ (follows from the normal equations).
- The prediction errors are then

$$X_h - \hat{X}_h = Z_h$$

$$X_0 - \hat{X}_0 = (1 - \phi_1^2)X_0 - \phi_1 Z_1.$$

These are uncorrelated for $h \geq 2$. So $\kappa(h) = 0$.

- The argument can be generalized to the case of an AR(p), for which $\kappa(h) = 0$ when $h > p$.

Proposition 6.5.5.

If $\{X_t\}$ has mean zero, the PACF at lag h is given by solving the Yule-Walker equations of order h , and taking the last coefficient, i.e., letting \underline{e}_h denote the length h unit vector with one in the last position,

$$\kappa(h) = \underline{e}_h' \Gamma_h^{-1} \underline{\gamma}_h.$$

Exercise 6.54. PACF of MA(q)

- We use the formula of Proposition 6.5.5 to compute the PACF for the MA(3) process with $\theta(z) = 1 + .4z + .2z^2 - .3z^3$.
- First we load the ARMAauto.r function from earlier notebooks.

```

polymult <- function(a,b) {
  bb <- c(b,rep(0,length(a)-1))
  B <- toeplitz(bb)
  B[lower.tri(B)] <- 0
  aa <- rev(c(a,rep(0,length(b)-1)))
  prod <- B %*% matrix(aa,length(aa),1)
  return(rev(prod[,1]))
}

ARMAauto <- function(phi,theta,maxlag)
{
  p <- length(phi)
  q <- length(theta)
  gamMA <- polymult(c(1,theta),rev(c(1,theta)))
  gamMA <- gamMA[(q+1):(2*q+1)]
  if (p > 0)
  {
    Amat <- matrix(0,nrow=(p+1),ncol=(2*p+1))
    for(i in 1:(p+1))
    {
      Amat[i,i:(i+p)] <- c(-1*rev(phi),1)
    }
    Amat <- cbind(Amat[, (p+1)],as.matrix(Amat[, (p+2):(2*p+1)] +
      t(matrix(apply(t(matrix(Amat[,1:p],p+1,p)),2,rev),p,p+1)))
    Bmat <- matrix(0,nrow=(q+1),ncol=(p+q+1))
    for(i in 1:(q+1))
    {
      Bmat[i,i:(i+p)] <- c(-1*rev(phi),1)
    }
    Bmat <- t(matrix(apply(t(Bmat),2,rev),p+q+1,q+1))
    Bmat <- matrix(apply(Bmat,2,rev),q+1,p+q+1)
    Bmat <- Bmat[,1:(q+1)]
    Binv <- solve(Bmat)
    gamMix <- Binv %*% gamMA
    if (p <= q) { gamMix <- matrix(gamMix[1:(p+1),],p+1,1)
      } else gamMix <- matrix(c(gamMix,rep(0,(p-q))),p+1,1)
    gamARMA <- solve(Amat) %*% gamMix
  } else gamARMA <- gamMA[1]

  gamMA <- as.vector(gamMA)
  if (maxlag <= q) gamMA <- gamMA[1:(maxlag+1)] else gamMA <- c(gamMA,rep(0,(maxlag-q)))
  gamARMA <- as.vector(gamARMA)
  if (maxlag <= p) gamARMA <- gamARMA[1:(maxlag+1)] else {
    for(k in 1:(maxlag-p))
    {
      len <- length(gamARMA)
      acf <- gamMA[p+1+k]
      if (p > 0) acf <- acf + sum(phi*rev(gamARMA[(len-p+1):len]))
      gamARMA <- c(gamARMA,acf)
    }
  }
  return(gamARMA)
}

```

- Then we implement Proposition 6.5.5.

```

armapq.pacf <- function(ar.coefs,ma.coefs,max.lag)
{
  gamma <- ARMAauto(ar.coefs,ma.coefs,max.lag)
  kappa <- NULL
  for(k in 1:max.lag)
  {
    new.kappa <- solve(toeplitz(gamma[1:k]),gamma[2:(k+1)])[k]
    kappa <- c(kappa,new.kappa)
  }
  return(kappa)
}


```

- Then we apply to the given MA(3) process.

```

ma.coefs <- c(.4,.2,-.3)
kappa <- armapq.pacf(NULL,ma.coefs,10)
plot(ts(kappa,start=1),xlab="Lag",ylab="Partial Autocorrelation",
      ylim=c(min(kappa),max(kappa)),type="h")

```



Exercise 6.55. PACF pf ARMA(p,q)


- Compute the PACF of Example 5.5.7, which is an ARMA(1,2) with $\phi(z) = 1 - .5z$ and $\theta(z) = 1 + (5/6)z + (1/6)z^2$.

```

phi1 <- .5
theta1 <- 5/6
theta2 <- 1/6

```

```
kappa <- armapq.pacf(phi1,c(theta1,theta2),10)
plot(ts(kappa,start=1),xlab="Lag",ylab="Partial Autocorrelation",
     ylim=c(min(kappa),max(kappa)),type="h")
```



Lesson 6-6: AR and MA Identification

- How do we determine a model to be fitted? AR, MA, ARMA, or something else?
- How do we determine the order of the model?

Paradigm 6.6.1. Characterizing AR and MA Processes

- The autocorrelation function (ACF), inverse autocorrelation (IACF), and partial autocorrelation function (PACF) have distinctive behavior for AR and MA processes.

ACF

- For an $MA(q)$ process, the ACF truncates at lag q , i.e., $\gamma(h) = 0$ if $|h| > q$. However, it is possible for $\gamma(h) = 0$ for $0 < h < q$ as well.
- For an $AR(p)$ process, or for an $ARMA(p,q)$ process (with $p > 0$), the ACF decays at geometric rate. The correlations can oscillate, but they are bounded by some $Cr^{|h|}$ for $0 < r < 1$ and $C > 0$.

IACF

- Generalize Exercise 6.34 to get IACF behavior for $AR(p)$ processes.
- For an $AR(p)$ process, the IACF truncates at lag p , i.e., $\zeta(h) = 0$ if $|h| > p$.
- for an $MA(q)$ process, or for an $ARMA(p,q)$ process (with $q > 0$), the IACF decays at geometric rate.

PACF

- For an $AR(p)$ process, the PACF truncates at lag p , i.e., $\kappa(h) = 0$ if $|h| > p$.
- for an $MA(q)$ process, or for an $ARMA(p,q)$ process (with $q > 0$), the PACF decays at geometric rate.

Finding Truncation

We can plot estimates of the ACF and PACF, and see if there is a lag cut-off where one or the other seems to negligible.

Example 6.6.2. MA(3) Identification

- Suppose we observe the ACF, IACF, and PACF of a process.

```
polymult <- function(a,b) {
  bb <- c(b,rep(0,length(a)-1))
  B <- toeplitz(bb)
  B[lower.tri(B)] <- 0
  aa <- rev(c(a,rep(0,length(b)-1)))
  prod <- B %*% matrix(aa,length(aa),1)
  return(rev(prod[,1]))
}

ARMAauto <- function(phi,theta,maxlag)
{
  p <- length(phi)
  q <- length(theta)
  gamMA <- polymult(c(1,theta),rev(c(1,theta)))
  gamMA <- gamMA[(q+1):(2*q+1)]
  if (p > 0)
  {
    Amat <- matrix(0,nrow=(p+1),ncol=(2*p+1))
    for(i in 1:(p+1))
    {
      Amat[i,i:(i+p)] <- c(-1*rev(phi),1)
    }
    Amat <- cbind(Amat[, (p+1)],as.matrix(Amat[, (p+2):(2*p+1)] +
      t(matrix(apply(t(matrix(Amat[,1:p],p+1,p)),2,rev),p,p+1)))
    Bmat <- matrix(0,nrow=(q+1),ncol=(p+q+1))
    for(i in 1:(q+1))
    {
      Bmat[i,i:(i+p)] <- c(-1*rev(phi),1)
    }
    Bmat <- t(matrix(apply(t(Bmat),2,rev),p+q+1,q+1))
    Bmat <- matrix(apply(Bmat,2,rev),q+1,p+q+1)
    Bmat <- Bmat[,1:(q+1)]
    Binv <- solve(Bmat)
    gamMix <- Binv %*% gamMA
    if (p <= q) { gamMix <- matrix(gamMix[1:(p+1),],p+1,1)
      } else gamMix <- matrix(c(gamMix,rep(0,(p-q))),p+1,1)
    gamARMA <- solve(Amat) %*% gamMix
  } else gamARMA <- gamMA[1]

  gamMA <- as.vector(gamMA)
  if (maxlag <= q) gamMA <- gamMA[1:(maxlag+1)] else gamMA <- c(gamMA,rep(0,(maxlag-q)))
}
```

```

gamARMA <- as.vector(gamARMA)
if (maxlag <= p) gamARMA <- gamARMA[1:(maxlag+1)] else {
  for(k in 1:(maxlag-p))
  {
    len <- length(gamARMA)
    acf <- gamMA[p+1+k]
    if (p > 0) acf <- acf + sum(phi*rev(gamARMA[(len-p+1):len]))
    gamARMA <- c(gamARMA,acf)
  } }
return(gamARMA)
}

armapq.pacf <- function(ar.coefs,ma.coefs,max.lag)
{
  gamma <- ARMAauto(ar.coefs,ma.coefs,max.lag)
  kappa <- NULL
  for(k in 1:max.lag)
  {
    new.kappa <- solve(toeplitz(gamma[1:k]),gamma[2:(k+1)])[k]
    kappa <- c(kappa,new.kappa)
  }
  return(kappa)
}

```

- We construct and plot these functions for an MA(3) process.

```

ma.coefs <- c(.4,.2,-.3)
gamma <- ARMAauto(NULL,ma.coefs,10)
rho <- gamma/gamma[1]
xi <- ARMAauto(-1*ma.coefs,NULL,10)
zeta <- xi/xi[1]
kappa <- armapq.pacf(NULL,ma.coefs,10)


```

- The ACF plot. We start at lag 1, since the lag 0 value is always zero.

```


plot(ts(rho[-1],start=1),xlab="Lag",ylab="Autocorrelation",
      ylim=c(-1,1),type="h")
abline(h=0,col=2)

```




- The IACF plot. We start at lag 1, since the lag 0 value is always zero.

```
plot(ts(zeta[-1],start=1),xlab="Lag",ylab="Inverse Autocorrelation",  
     ylim=c(-1,1),type="h")  
abline(h=0,col=2)
```



- The PACF plot. We start at lag 1, since the lag 0 value is not defined.

```
plot(ts(kappa,start=1),xlab="Lag",ylab="Partial Autocorrelation",
     ylim=c(-1,1),type="h")
abline(h=0,col=2)
```




Example 6.6.3. AR(4) Identification

- Suppose we observe the ACF, IACF, and PACF of a process.
- We construct and plot these functions for an AR(4) process.

```
ar.coefs <- c(.8, -.3, .2, .1)
gamma <- ARMAauto(ar.coefs, NULL, 10)
rho <- gamma/gamma[1]
xi <- ARMAauto(NULL, -1*ar.coefs, 10)
zeta <- xi/xi[1]
kappa <- armapq.pacf(ar.coefs, NULL, 10)
```


- The ACF plot. We start at lag 1, since the lag 0 value is always zero.

```
plot(ts(rho[-1], start=1), xlab="Lag", ylab="Autocorrelation",
     ylim=c(-1,1), type="h")
abline(h=0, col=2)
```




- The IACF plot. We start at lag 1, since the lag 0 value is always zero.

```
plot(ts(zeta[-1],start=1),xlab="Lag",ylab="Inverse Autocorrelation",  
     ylim=c(-1,1),type="h")  
abline(h=0,col=2)
```



- The PACF plot. We start at lag 1, since the lag 0 value is not defined.

```
plot(ts(kappa,start=1),xlab="Lag",ylab="Partial Autocorrelation",
     ylim=c(-1,1),type="h")
abline(h=0,col=2)
```



Paradigm 6.6.7. Identification by Whitening

- Suppose we apply some filter $\psi(B)$ to the data $\{X_t\}$, and the output appears to be white noise $\{Z_t\}$ (e.g., we ran some statistical tests of serial independence).
- $\psi(B)$ is called a *whitening filter*.
- We infer that $X_t = \psi(B)^{-1}Z_t$, which gives a model for $\{X_t\}$.
- So we can try out classes of filters $\psi(B)$, attempt to whiten the data, and deduce the original model.

Example 6.6.8. $AR(p)$ Whitening Models


- Consider the class of filters $\psi(B) = 1 - \sum_{j=1}^p \psi_j B^j$, which are AR polynomial filters.
- We would apply these to the data, seeking p and coefficient values such that the data is whitened.
- We can estimate coefficients using ordinary least squares (or the Yule-Walker method, discussed later), for any p . These are fast to calculate, so we can just try over many choices of p .

Exercise 6.6.1. Whitening an $AR(p)$ Process

- We implement the method of Example 6.6.8, and apply to an $AR(2)$ simulation.

```
arp.sim <- function(n,burn,ar.coefs,innovar)
{
  p <- length(ar.coefs)
  z <- rnorm(n+burn+p,sd=sqrt(innovar))
  x <- z[1:p]
  for(t in (p+1):(p+n+burn))
  {
```


- ```
set.seed(777)
```



- ```
covars <- as.matrix(x.sim[-n])
```

```

ar.fit <- lm(x.sim[-seq(1,p)] ~ covars - 1)
coeffs[[p]] <- ar.fit$coefficients
resids[[p]] <- ar.fit$residuals
covars <- cbind(covars[-1,],x.sim[-seq(n-p,n)])
}


```

- We plot the filter outputs (the regression residuals) and estimates of the ACF.


```


for(p in 1:5)
{
  plot.ts(resids[[p]],ylab="")
  acf(resids[[p]])
}

```





Series residu[[p]]







Series resid_s[[p]]







Series residu[[p]]






Series residu[[p]]







- The correct filter is for $p = 2$. We print the coefficients and their estimates.

```
print(c(phi1,phi2))
```

```
## [1] 1.385641 -0.640000
```

```
print(coeffs[[2]])
```

```
## covars1 covars2
```

```
## 1.4773729 -0.6966449
```

Lesson 7-2: Discrete Fourier Transform

We introduce the data analysis tool of Discrete Fourier Transform.

Definition 7.2.3.

Given a sample X_1, \dots, X_n , the **Discrete Fourier Transform** (DFT) at the Fourier frequency $\lambda_l = 2\pi l/n$ (for $[n/2] - n + 1 \leq l \leq [n/2]$) is

$$\tilde{X}(\lambda_l) = n^{-1/2} \sum_{t=1}^n X_t e^{-i\lambda_l t}.$$

Definition 7.2.4.

- The **periodogram** $I(\lambda)$ is a non-negative function of $\lambda \in [-\pi, \pi]$, constructed from the sample X_1, \dots, X_n :

$$I(\lambda) = n^{-1} \left| \sum_{t=1}^n (X_t - \bar{X}) e^{-i\lambda t} \right|^2.$$

- So $I(0) = 0$. Sometimes we consider an “uncentered” periodogram, where there is no centering by the sample mean. We denote this by $\tilde{I}(\lambda)$, and

$$\tilde{I}(\lambda_l) = |\tilde{X}(\lambda_l)|^2.$$

- The periodogram is an empirical version of the spectral density, and shares certain properties. Higher values correspond to cyclical effects in the data.

Proposition 7.2.7.

- Let the vector of DFTs be denoted $\tilde{\underline{X}}$, which has components $\tilde{X}(\lambda_l)$.
- This is a linear function of the sample vector \underline{X} :

$$\tilde{\underline{X}} = Q^* \underline{X},$$

where Q was defined in Definition 6.4.3.

- Since Q is unitary, $Q^{-1} = Q^*$, so we can recover the data from the DFT vector via

$$\underline{X} = Q \tilde{\underline{X}}.$$

Exercise 7.14. DFT of an AR(1).

- We simulate an AR(1).

```
arp.sim <- function(n,burn,ar.coefs,innovar)
{
  p <- length(ar.coefs)
  z <- rnorm(n+burn+p,sd=sqrt(innovar))
  x <- z[1:p]
  for(t in (p+1):(p+n+burn))
  {
    next.x <- sum(ar.coefs*x[(t-1):(t-p)]) + z[t]
    x <- c(x,next.x)
  }
  x <- x[(p+burn+1):(p+burn+n)]
  return(x)
}

phi <- .8
innovar <- 1
n <- 200
x.sim <- arp.sim(n,500,phi,innovar)
```

- We compute the DFT. We begin with code to compute Q .

```
get.qmat <- function(mesh)
{
  mesh2 <- floor(mesh/2)
  inds <- seq(mesh2-mesh+1,mesh2)
```

```

Q.mat <- exp(1i*2*pi*mesh^{-1}*t(t(seq(1,mesh)) %% inds))*mesh^{-1/2}
return(Q.mat)
}
Q.mat <- get.qmat(n)


```

- Then we use Proposition 7.2.7 to get the DFT vector. We plot the real part, imaginary part, and the modulus.

```

x.dft <- Conj(t(Q.mat)) %*% x.sim
plot(ts(Re(x.dft),start=-floor(n/2)+1,frequency=1),
     xlab="Frequency Index",ylab="Real Part")


```




```

plot(ts(Im(x.dft),start=-floor(n/2)+1,frequency=1),
     xlab="Frequency Index",ylab="Imaginary Part")

```



```
plot(ts(Mod(x.dft)^2,start=-floor(n/2)+1,frequency=1),  
      xlab="Frequency Index",ylab="Squared Modulus")
```



Corollary 7.2.9. Decorrelation Property of the DFT.


- Let X_1, \dots, X_n be a sample from a mean zero, covariance stationary time series with absolutely summable ACVF and spectral density f . Then $\tilde{\underline{X}}$ has approximate covariance matrix $\text{diag}\{f(\lambda_{[n/2]-n+k})\}$.
- This follows from Theorem 6.4.5:

$$\text{Cov}[\tilde{\underline{X}}] = Q^* \text{Cov}[\underline{X}] Q \approx \Lambda.$$


Exercise 7.18. DFT of Wolfer Sunspots.

- We compute the DFT of the Wolfer sunspot data, and plot.


```
wolfer <- read.table("wolfer.dat")
wolfer <- ts(wolfer, start=1749, frequency=12)
n <- length(wolfer)
Q.mat <- get.qmat(n)
x.dft <- Conj(t(Q.mat)) %*% wolfer
plot(ts(Re(x.dft), start=-floor(n/2)+1, frequency=1),
     xlab="Frequency Index", ylab="Real Part")
```



```
plot(ts(Im(x.dft),start=-floor(n/2)+1,frequency=1),  
     xlab="Frequency Index",ylab="Imaginary Part")
```

```
plot(ts(log(Mod(x.dft)^2),start=-floor(n/2)+1,frequency=1),  
     xlab="Frequency Index",ylab="Log Squared Modulus")
```



- We see higher values of the squared modulus (the uncentered periodogram) near frequency zero.


Exercise 7.20. DFT of Mauna Loa Growth Rate.

- We compute the DFT of the Mauna Loa growth rate, and plot.

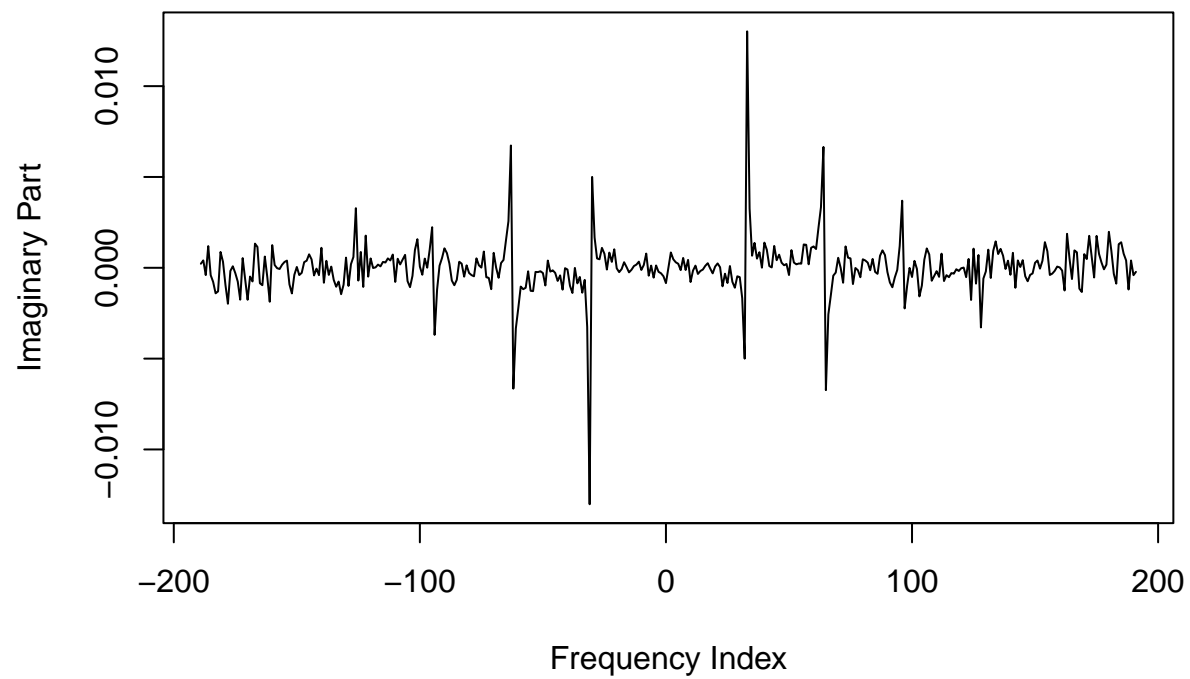
```

mau <- read.table("mauna.dat",header=TRUE,sep="")
mau <- ts(mau,start=1958,frequency=12)
mau.gr <- diff(log(mau))
n <- length(mau.gr)
Q.mat <- get.qmat(n)
x.dft <- Conj(t(Q.mat)) %*% mau.gr
plot(ts(Re(x.dft),start=-floor(n/2)+1,frequency=1),
     xlab="Frequency Index",ylab="Real Part")


```



```
plot(ts(Im(x.dft),start=-floor(n/2)+1,frequency=1),  
     xlab="Frequency Index",ylab="Imaginary Part")
```



```
plot(ts(log(Mod(x.dft)^2),start=-floor(n/2)+1,frequency=1),  
     xlab="Frequency Index",ylab="Log Squared Modulus")
```



- We see higher values of the uncentered periodogram in the shape of peaks, at four non-zero frequencies. These correspond to cyclical seasonal effects.

Lesson 7-3: Spectral Representation

We discuss a representation of a stationary time series as a sum of stochastic cosines.

Definition 7.3.1.

- For a given spectral distribution F , a **spectral increment process** Z is a complex-valued continuous-time stochastic process defined on the interval $[-\pi, \pi]$, which has mean zero with *orthogonal increments*: for $\lambda_1 < \lambda_2 < \lambda_3 < \lambda_4$, the random variables $Z(\lambda_2) - Z(\lambda_1)$ and $Z(\lambda_4) - Z(\lambda_3)$ are orthogonal. Also,

$$\text{Var}[Z(\lambda_2) - Z(\lambda_1)] = \frac{1}{2\pi} (F(\lambda_2) - F(\lambda_1)).$$

- The variance of a complex random variable is the expectation of its squared modulus.
- We abbreviate the above expression by

$$\text{Var}[dZ(\lambda)] = \frac{1}{2\pi} dF(\lambda).$$

Paradigm 7.3.4. Time Series Defined as a Stochastic Integral

- We can define a time series via a stochastic integral as follows:

$$X_t = \int_{-\pi}^{\pi} e^{i\lambda t} dZ(\lambda).$$

- This resembles a sum of stochastic cosines, where $dZ(\lambda)$ is the amplitude for a sinusoid $e^{i\lambda t}$.
- Then $\{X_t\}$ has mean zero and autocovariance

$$\text{Cov}[X_{t+h}, X_t] = \int_{-\pi}^{\pi} \int_{-\pi}^{\pi} e^{i\lambda(t+h)} e^{-i\omega t} \text{Cov}[dZ(\lambda), dZ(\omega)] = \frac{1}{2\pi} \int_{-\pi}^{\pi} e^{i\lambda h} dF(\lambda).$$

- So $\{X_t\}$ is weakly stationary with ACVF $\gamma(h)$ given by above formula.
- Conversely: any mean zero weakly stationary time series $\{X_t\}$ with spectral distribution function F can be represented by the above stochastic integral!

Corollary 7.3.8.

- Suppose $\{X_t\}$ is a mean zero weakly stationary time series with spectral representation, and let $Y_t = \psi(B)X_t$. Then

$$Y_t = \sum_j \psi_j X_{t-j} = \sum_j \psi_j \int_{-\pi}^{\pi} e^{i\lambda(t-j)} dZ(\lambda) = \int_{-\pi}^{\pi} \sum_j \psi_j e^{-i\lambda j} e^{i\lambda t} dZ(\lambda) = \int_{-\pi}^{\pi} \psi(e^{-i\lambda}) e^{i\lambda t} dZ(\lambda).$$

- So $\{Y_t\}$ also has a spectral representation, but its increment process is $\psi(e^{-i\lambda})dZ(\lambda)$, where $\psi(e^{-i\lambda})$ is the filter frequency response function. We can use this result to understand the impact of a filter in the frequency domain.
- The ACVF is

$$\text{Cov}[Y_{t+h}, Y_t] = \frac{1}{2\pi} \int_{-\pi}^{\pi} e^{i\lambda h} |\psi(e^{-i\lambda})|^2 dF(\lambda).$$

Example 7.3.10. Time Shift

- Consider $\psi(B) = B^k$, a shift by k time units. Then $\psi(e^{-i\lambda}) = e^{-i\lambda k}$ and $Y_t = \int_{-\pi}^{\pi} e^{i\lambda(t-k)} dZ(\lambda)$.
- This is the same as $B^k X_t = X_{t-k}$.

Definition 7.3.11.

- We can decompose the frequency response function into the **gain** function and the **phase delay** function, each with an interpretation from the spectral representation.
- We use the polar decomposition of a complex number: for any λ ,

$$\psi(e^{-i\lambda}) = |\psi(e^{-i\lambda})| \exp\{i \text{Arg} \psi(e^{-i\lambda})\}.$$

- The magnitude $|\psi(e^{-i\lambda})|$ is called the **gain** function of the filter. It is computed by taking the square root of sum of squares of real and imaginary parts.
- The angular portion $\text{Arg} \psi(e^{-i\lambda})$ is called the **phase** function of the filter. It is computed by taking the arc tangent of the ratio of imaginary to real parts.
- When the phase function is differentiable with respect to λ , we define the **phase delay** via

$$\Upsilon(\lambda) = \frac{-\text{Arg} \psi(e^{-i\lambda})}{\lambda}$$

for $\lambda \neq 0$, and the limit of such for $\lambda = 0$.

- The phase delay may be discontinuous in λ .
- The gain is an even function; both phase and phase delay are odd. So we usually just plot them over $[0, \pi]$ instead of $[-\pi, \pi]$.

Fact 7.3.12. Action of Phase Delay.

- From Corollary 7.3.8,

$$Y_t = \int_{-\pi}^{\pi} e^{i(\lambda t + \text{Arg} \psi(e^{-i\lambda}))} |\psi(e^{-i\lambda})| dZ(\lambda) = \int_{-\pi}^{\pi} e^{i\lambda(t - \Upsilon(\lambda))} |\psi(e^{-i\lambda})| dZ(\lambda).$$

- So at frequency λ , time t is delayed by $\Upsilon(\lambda)$ time units.
- Also, the gain modifies the autocovariances.

Example 7.3.13. Simple Moving Average Filters Cause Delay

- Consider the simple moving average filter $\psi(B) = (1 + B + B^2)/3$, which gives the average of the past and present 3 observations.
- We directly compute the frequency response function:

$$\psi(e^{-i\lambda}) = \frac{1 + e^{-i\lambda} + e^{-i2\lambda}}{3} = e^{-i\lambda} \frac{e^{i\lambda} + 1 + e^{-i\lambda}}{3} = e^{-i\lambda} \frac{1 + 2\cos(\lambda)}{3}.$$

- The gain function is

$$\frac{1}{3}|1 + 2\cos(\lambda)|.$$

- Note that $1 + 2\cos(\lambda)$ is non-negative for $\lambda \in [0, 2\pi/3]$, so the phase function equals $-\lambda$ over that set. Otherwise, we need a -1 factor which leads to a phase function equal to $-\lambda - \pi$.
- The phase delay function is then

$$\Upsilon(\lambda) = \begin{cases} 1 & \text{if } \lambda \in [0, 2\pi/3] \\ 1 + \pi/\lambda & \text{else.} \end{cases}$$

- Since this function is positive, the filter always provides a delay.
- The gain function attenuates higher frequencies, due to the cosine shape, so the filter is a “low-pass.”

Example 7.3.14. Differencing Causes an Advance

- Consider the differencing filter $\psi(B) = 1 - B$.
- The frequency response function is

$$\psi(e^{-i\lambda}) = 1 - e^{-i\lambda} = 1 - \cos(\lambda) - i\sin(\lambda).$$

- The squared gain function is $2 - 2\cos(\lambda)$.
- The phase function is $(\pi - \lambda)/2$.
- Away from $\lambda = 0$, the phase delay is $\Upsilon(\lambda) = .5(1 - \pi/\lambda)$. This is negative for $\lambda \in (0, \pi]$, so differencing causes an advance.
- The gain function attenuates lower frequencies, so the filter is a “high-pass.”


Exercise 7.29. Phase and Gain for Simple Moving Average.

- Consider the simple moving average of order p : $\psi(B) = (2p + 1)^{-1} \sum_{j=-p}^p B^j$.
- It can be shown that


$$\psi(e^{-i\lambda}) = \frac{\sin(\lambda(p + 1/2))}{(2p + 1)\sin(\lambda/2)}$$

- We use this formula to encode and display the gain and phase functions, as well as the real and imaginary parts of the frequency response function.


```
lambda <- pi*seq(-1000,1000)/1000
p <- 3
simplema.frf <- sin((p+1/2)*lambda)/((2*p+1)*sin(lambda/2))
simplema.gain <- Mod(simplema.frf)
simplema.phase <- atan(Im(simplema.frf)/Re(simplema.frf))
simplema.delay <- -1*simplema.phase/lambda
plot(ts(Re(simplema.frf),start=-1,frequency=1000),
     xlab="Frequency",ylab="Real Part Frf")
```




```
plot(ts(Im(simplema.frf),start=-1,frequency=1000),  
     xlab="Frequency",ylab="Imaginary Part Frf")
```


```
plot(ts(simplema.gain,start=-1,frequency=1000),  
     xlab="Frequency",ylab="Gain")
```



```
plot(ts(simplema.phase,start=-1,frequency=1000),  
     xlab="Frequency",ylab="Phase")
```




```
plot(ts(simplema.delay,start=-1,frequency=1000),  
     xlab="Frequency",ylab="Phase Delay")
```




Exercise 7.30. Phase and Gain for Seasonal Aggregation Filter.

- Consider the seasonal aggregation filter: $\psi(B) = \sum_{j=0}^{s-1} B^j$.
- We encode and display the gain and phase functions, as well as the real and imaginary parts of the frequency response function.


```
lambda <- pi*seq(-1000,1000)/1000
s <- 12
seasagg.frf <- exp(-1i*lambda*(s-1)/2)*sin((s/2)*lambda)/sin(lambda/2)
seasagg.gain <- Mod(seasagg.frf)
seasagg.phase <- atan(Im(seasagg.frf)/Re(seasagg.frf))
seasagg.delay <- -1*seasagg.phase/lambda
seasagg.delay[1001] <- NA
plot(ts(Re(seasagg.frf),start=-1,frequency=1000),
     xlab="Frequency",ylab="Real Part Frf")
```




```
plot(ts(Im(seasagg.frf),start=-1,frequency=1000),  
     xlab="Frequency",ylab="Imaginary Part Frf")
```



```
plot(ts(seasagg.gain,start=-1,frequency=1000),  
     xlab="Frequency",ylab="Gain")
```



```
plot(ts(seasagg.phase,start=-1,frequency=1000),  
     xlab="Frequency",ylab="Phase")
```



```
plot(ts(seasagg.delay,start=-1,frequency=1000),  
     xlab="Frequency",ylab="Phase Delay")
```