

Time Series: A First Course with Bootstrap Starter

Contents

Lesson 11-1: Nonlinear Processes	2
Paradigm 11.1.1. Linear, Causal, and Gaussian Processes	2
Example 11.1.2. Linear but Non-Gaussian $AR(p)$	2
Example 11.1.5. Nonlinear Autoregression	2
Example 11.1.6. Autoregressive Conditional Heteroscedasticity	2
Example 11.1.7. Prediction of Lognormal Time Series	3
Exercise 11.3. Lognormal Time Series Prediction Error	3
Lesson 11-3: ARCH Process	4
Remark 11.3.1. Volatility Clustering	4
Proposition 11.3.2.	5
Theorem 11.3.6.	5
Remark 11.3.8. The Market Efficiency Axiom	5
Paradigm 11.3.9. Fitting an $ARCH(p)$ Model	5
Example 11.3.10. Fitting a Gaussian $ARCH(p)$ to Dow Log Returns	5
Lesson 11-4: GARCH Processes	7
Definition 11.4.1.	7
Theorem 11.4.3.	7
Example 11.4.4. $GARCH(1,1)$	7
Corollary 11.4.5.	8
Remark 11.4.6. Fitting a GARCH Model	8
Example 11.4.7. Fitting a Gaussian $GARCH(1,1)$ to Dow Log Returns	8
Example 11.4.8. Fat-Tailed $GARCH(1,1)$ for Dow Log Returns	10
Lesson 11-5: Bi-spectral Density	11
Definition 11.5.1.	11
Remark 11.5.2. Cumulant Functions	12
Fact 11.5.3. Symmetries of the Third Cumulant	12
Definition 11.5.6.	12
Example 11.5.7. Bi-spectral Density of an $ARCH(1)$ Process	12
Theorem 11.5.8.	12
Lesson 11-6: Volatility Filtering	12
Definition 11.6.2	12
Paradigm 11.6.3. Unbiased Volatility Filtering	13
Example 11.6.4. Volatility Filtering of Dow Jones Log Returns	13
Definition 11.6.5.	13
Paradigm 11.6.6. NoVaS with Simple Weights	14
Example 11.6.7. Simple NoVaS Applied to Dow Log Returns	14

Lesson 11-1: Nonlinear Processes

- We discuss various types of nonlinear processes.

Paradigm 11.1.1. Linear, Causal, and Gaussian Processes

- $\{X_t\}$ is Gaussian if all its finite-dimensional marginal distributions are multivariate Gaussian.
- The best one-step ahead predictor for a Gaussian time series is a *linear* function of the sample.
- More broadly, a time series is *linear* if it can be written

$$X_t = \sum_{j=-\infty}^{\infty} \psi_j \epsilon_{t-j}$$

for $\{\epsilon_t\}$ i.i.d.

- Some linear time series are *causal*, where $\psi(z)$ is a power series.
- Some causal linear time series are *invertible*, where roots of $\psi(z)$ are outside the unit circle.
- A stationary Gaussian time series with positive spectral density is linear, causal, and invertible.

Example 11.1.2. Linear but Non-Gaussian AR(p)

- Consider a non-Gaussian AR(p) process:

$$X_t = \sum_{j=1}^p \pi_j X_{t-j} + \epsilon_t$$

for $\{\epsilon_t\}$ i.i.d. (but not Gaussian).

- The best one-step ahead predictor is still linear: supposing that $t > p$,

$$\mathbb{E}[X_t | \underline{X}_{t-1}] = \sum_{j=1}^p \pi_j \mathbb{E}[X_{t-j} | \underline{X}_{t-1}] + \mathbb{E}[\epsilon_t | \underline{X}_{t-1}] = \sum_{j=1}^p \pi_j X_{t-j}.$$

Example 11.1.5. Nonlinear Autoregression

- We can generalize the regression in Example 11.1.2 to be a nonlinear function of past data.
- An order one nonlinear autoregression has the form

$$X_t = g(X_{t-1}) + Z_t,$$

where $\{Z_t\}$ is i.i.d.

- E.g., $g(x) = \exp(x)$.
- The best one-step ahead predictor is

$$\mathbb{E}[X_t | \underline{X}_{t-1}] = g(X_{t-1}).$$

Example 11.1.6. Autoregressive Conditional Heteroscedasticity

- $\{X_t\}$ is an AutoRegressive Conditional Heteroscedastic (ARCH) process of order p if

$$X_t = \sigma_t Z_t$$
$$\sigma_t^2 = a_0 + \sum_{j=1}^p a_j X_{t-j}^2,$$

where $\{Z_t\}$ is i.i.d. $(0, 1)$, and $a_j \geq 0$.

- Then Z_t is independent of past values of the process, and $\{X_t^2\}$ is a nonlinear process whose best predictor is linear (shown in Theorem 11.3.6 below).

Example 11.1.7. Prediction of Lognormal Time Series

- The lognormal times series $\{X_t\}$ is defined by $X_t = \exp Y_t$, where $\{Y_t\}$ is Gaussian.
- Say $\{Y_t\}$ is mean zero with ACVF $\gamma_Y(h)$. Then $\mathbb{E}[X_t] = \exp\{\gamma_Y(0)/2\}$ and

$$\gamma_X(h) = \exp\{\gamma_Y(0)\} (\exp\{\gamma_Y(h)\} - 1).$$

- Suppose that σ^2 is the asymptotic prediction variance of Y_t , and \hat{Y}_{t+1} is the best predictor of Y_{t+1} . Then

$$\hat{X}_{t+1} = \exp\{\hat{Y}_{t+1} + \sigma^2/2\},$$

and the nonlinear prediction error variance is

$$\exp\{2\gamma_Y(0)\} \cdot (1 - \exp\{-\sigma^2\}).$$

Exercise 11.3. Lognormal Time Series Prediction Error

- We numerically compute the linear and nonlinear prediction error variances.
- The spectral density of $\{X_t\}$ is

$$f(\lambda) = \exp\{\gamma_Y(0)\} \sum_{h=-\infty}^{\infty} (\exp\{\gamma_Y(h)\} - 1) e^{-i\lambda h}.$$

- So the best linear prediction error variance is $\exp\{\langle \log f \rangle_0\}$.
- We consider the case where $\{Y_t\}$ is a Gaussian AR(1) of parameter $\phi = .8$ and $\sigma^2 = 1$.

```
polymul <- function(a,b)
{
  bb <- c(b,rep(0,length(a)-1))
  B <- toeplitz(bb)
  B[lower.tri(B)] <- 0
  aa <- rev(c(a,rep(0,length(b)-1)))
  prod <- B %%% matrix(aa,length(aa),1)
  return(rev(prod[,1]))
}

ARMAauto <- function(phi,theta,maxlag)
{
  p <- length(phi)
  q <- length(theta)
  gamMA <- polymul(c(1,theta),rev(c(1,theta)))
  gamMA <- gamMA[(q+1):(2*q+1)]
  if (p > 0)
  {
    Amat <- matrix(0,nrow=(p+1),ncol=(2*p+1))
    for(i in 1:(p+1))
    {
      Amat[i,i:(i+p)] <- c(-1*rev(phi),1)
    }
    Amat <- cbind(Amat[, (p+1)], as.matrix(Amat[, (p+2):(2*p+1)] +
      t(matrix(apply(t(matrix(Amat[, 1:p], p+1, p)), 2, rev), p, p+1)))
    Bmat <- matrix(0,nrow=(q+1),ncol=(p+q+1))
    for(i in 1:(q+1))
    {
      Bmat[i,i:(i+p)] <- c(-1*rev(phi),1)
    }
  }
}
```

```

      Bmat <- t(matrix(apply(t(Bmat),2,rev),p+q+1,q+1))
      Bmat <- matrix(apply(Bmat,2,rev),q+1,p+q+1)
      Bmat <- Bmat[,1:(q+1)]
      Binv <- solve(Bmat)
      gamMix <- Binv %*% gamMA
      if (p <= q) { gamMix <- matrix(gamMix[1:(p+1),],p+1,1)
        } else gamMix <- matrix(c(gamMix,rep(0,(p-q))),p+1,1)
      gamARMA <- solve(Amat) %*% gamMix
    } else gamARMA <- gamMA[1]

    gamMA <- as.vector(gamMA)
    if (maxlag <= q) gamMA <- gamMA[1:(maxlag+1)] else gamMA <- c(gamMA,rep(0,(maxlag-q)))
    gamARMA <- as.vector(gamARMA)
    if (maxlag <= p) gamARMA <- gamARMA[1:(maxlag+1)] else {
    for(k in 1:(maxlag-p))
    {
      len <- length(gamARMA)
      acf <- gamMA[p+1+k]
      if (p > 0) acf <- acf + sum(phi*rev(gamARMA[(len-p+1):len]))
      gamARMA <- c(gamARMA,acf)
    } }
    return(gamARMA)
  }

phi <- .8
sigma <- 1
gamma <- ARMAauto(phi,NULL,100)*sigma^2
mesh <- 1000
lambda <- pi*seq(0,mesh)/mesh
f.spec <- (exp(gamma[1])-1)*cos(0*lambda)
for(h in 1:100)
{
  f.spec <- f.spec + 2*(exp(gamma[h+1])-1)*cos(h*lambda)
}
f.spec <- exp( gamma[1] ) * f.spec
lin.pred.mse <- exp( mean(log(f.spec)) )
nonlin.pred.mse <- exp(2*gamma[1])*(1- exp(-sigma^2))
print(c(lin.pred.mse,nonlin.pred.mse))

```

```
## [1] 170.0569 163.5110
```

- Notice that the nonlinear predictor has lower MSE than the linear predictor.

Lesson 11-3: ARCH Process

- For financial data, we need models that explain *fat tails* and *volatility clustering*.
- The ARCH model is a popular nonlinear time series model for financial data.

Remark 11.3.1. Volatility Clustering

- Financial time series exhibit periods of high variability followed by low variability.
- In intra-day trading data, volatility is higher at the beginning and end of the day.

Proposition 11.3.2.

- Let $\{X_t\}$ be an ARCH(p) process.
- It has mean zero, and conditional variance

$$\text{Var}[X_t|X_{t-1:}] = \sigma_t^2,$$

where $X_{t-1:}$ is shorthand for X_{t-1}, X_{t-2}, \dots

- The conditional variance σ_t^2 is called the *volatility*.

Theorem 11.3.6.

- Let $\{X_t\}$ be an ARCH(p) process such that $\delta = \sum_{j=1}^p a_j < 1$.
- Then $\{X_t\}$ is white noise with variance $a_0/(1 - \delta)$.
- If $\mathbb{E}[Z_t^4] < \infty$ and $\delta < \mathbb{E}[Z_t^4]^{-1/2}$, then $\{X_t^2\}$ is a causal AR(p) with respect to white noise inputs ϵ_t .
- Since the innovations ϵ_t are a dependent white noise, it can be shown that $\{X_t^2\}$ is a nonlinear process.

Remark 11.3.8. The Market Efficiency Axiom

- We can predict X_t^2 using its AR representation.
- But $\mathbb{E}[X_t|X_{t-1:}] = 0$.
- So we can say something about the magnitude of future returns, but not the direction.
- Knowing future values of log returns would permit an arbitrage opportunity.
- The *market efficiency axiom* states that there are no arbitrage opportunities in an efficient market.

Paradigm 11.3.9. Fitting an ARCH(p) Model

- We construct the likelihood of the ARCH model, based on the marginal distribution of Z_t .
- Given the past data x_{t-1}, x_{t-2}, \dots , we compute

$$s_t = \sqrt{a_0 + \sum_{j=1}^p a_j x_{t-j}^2}.$$

- Hence $p_Z(x_t/s_t)/s_t$ is the density for X_t given $X_{t-1:}$.
- The *pseudo-likelihood* is obtained by omitting some initial values. Its log is

$$\sum_{t=p+1}^n [\log p_Z(x_t/s_t) - \log s_t].$$

- This is a function of the parameters a_0, a_1, \dots, a_p , so we can numerically optimize.

Example 11.3.10. Fitting a Gaussian ARCH(p) to Dow Log Returns

- We illustrate the fitting of an ARCH model to the Dow log returns data.
- We first load the data and fit an AR(p) to the squares via Ordinary Least Squares, using AIC to select model order p .

```
dow <- read.table("dow.dat")
dow <- diff(log(dow[,1]))
dow <- ts(dow, start=c(2008,164), frequency=260)
dow.ols <- ar.ols(dow^2)
p <- dow.ols$order
print(p)
```

```
## [1] 33
```

- Then we utilize the pseudo-likelihood based on a Gaussian marginal distribution.
- This uses a reparameterization of the ARCH parameters, to ensure we get values in $[0, 1]$ that satisfy the constraint that $\delta < 1$.

```
psi2arch <- function(psi)
{
  p <- length(psi)-1
  a.0 <- exp(psi[1])
  if(p > 0)
  {
    r <- (1 + exp(-psi[2]))^(-1)
    if(p > 1)
    {
      a.1 <- (1 + sum(exp(-psi[3:(p+1)])))^(-1)
      a.j <- a.1
      for(j in 2:p)
      {
        a.j <- c(a.j,exp(-psi[j+1])*a.1)
      }
      a.j <- r*a.j
    } else { a.j <- r }
    a.0 <- c(a.0,a.j)
  }
  return(a.0)
}

lik.arch <- function(psi,data,df)
{
  p <- length(psi)-1
  T <- length(data)
  arch <- psi2arch(psi)

  lik <- 0
  for(t in (p+1):T)
  {
    arch.sd <- sqrt(arch[1] + sum(arch[-1]*data[(t-1):(t-p)]^2))
    if(df==Inf) { lik <- lik + (-2)*log(dnorm(data[t],sd=arch.sd)) } else {
      lik <- lik + (-2)*log(dt(data[t]/arch.sd,df=df)/arch.sd) }
  }
  return(lik)
}
```

- We initialize and run the optimization.

```
psi.init <- rep(0,p+1)
arch.fit <- optim(psi.init,lik.arch,data=dow,df=Inf,method="BFGS")
print(arch.fit)
```

```
## $par
## [1] -2033.2920745    63.3655168    -1.0359247    -0.9982537    -0.6378827
## [6]   -0.8151557   -0.8403049   -0.4084173   -0.3525442   -0.1679610
## [11]  -0.2513546    2.2654488    9.8935678    2.6956507    0.8310226
## [16]  10.4503981    6.3194832    2.5990406    1.7512673    9.4247571
## [21]  10.7428642    3.0778047    0.1038050    0.4003737   11.7614792
## [26]  11.8914237    2.1830580    9.0794927    0.3500038    0.2858920
```

```
## [31]      0.4896008      11.3106561      10.6978804      -0.2668674
##
## $value
## [1] -12517.6
##
## $counts
## function gradient
##      79      47
##
## $convergence
## [1] 0
##
## $message
## NULL

arch.par <- psi2arch(arch.fit$par)
print(round(arch.par,digits=3))

## [1] 0.000 0.041 0.116 0.111 0.078 0.093 0.095 0.062 0.058 0.049 0.053 0.004
## [13] 0.000 0.003 0.018 0.000 0.000 0.003 0.007 0.000 0.000 0.002 0.037 0.028
## [25] 0.000 0.000 0.005 0.000 0.029 0.031 0.025 0.000 0.000 0.054
```

Lesson 11-4: GARCH Processes

- We explore the GARCH process, a generalization of the ARCH process.

Definition 11.4.1.

- The *Generalized Autoregressive Conditionally Heteroscedastic* process of order p, q , or GARCH(p, q), is a process $\{X_t\}$ defined via

$$X_t = \sigma_t Z_t$$

$$\sigma_t^2 = a_0 + \sum_{j=1}^p a_j X_{t-j}^2 + \sum_{k=1}^q b_k \sigma_{t-k}^2.$$

- The $\{Z_t\}$ inputs are i.i.d. mean zero, variance 1.
- The parameters satisfy $a_j \geq 0$, $b_k \geq 0$.

Theorem 11.4.3.

- Let $\{X_t\}$ be a GARCH(p, q) process, and set $\delta = \sum_{j=1}^p a_j + \sum_{k=1}^q b_k$.
- If $\delta < 1$, then $\{X_t\}$ is strictly stationary with finite variance and mean zero.
- Also $\{X_t\}$ is a white noise with variance $a_0/(1 - \delta)$.
- If $\mathbb{E}[Z_t^4] < \infty$ and $\delta < \mathbb{E}[Z_t^4]^{-1/2}$, then $\{X_t^2\}$ is an ARMA(p^*, q) with respect to white noise inputs ϵ_t , and $p^* = \max\{p, q\}$.

Example 11.4.4. GARCH(1,1)

- Set $p = q = 1$ for the most popular GARCH specification.

$$\sigma_t^2 = a_0 + a_1 X_{t-1}^2 + b_1 \sigma_{t-1}^2.$$

- We can solve recursively for the volatility σ_t^2 :

$$\sigma_t^2 = \frac{a_0}{1 - b_1} + a_1 \sum_{j=0}^{\infty} b_1^j X_{t-1-j}^2.$$

- This is called the ARCH(∞) representation of the GARCH(1,1).

Corollary 11.4.5.

- Let $\{X_t\}$ be a GARCH(p,q) process with $\delta < 1$, such that $\mathbb{E}[Z_t^4] < \infty$ and $\delta < \mathbb{E}[Z_t^4]^{-1/2}$.
- Let $\theta(z) = 1 - \sum_{k=1}^q b_k z^k$ and $\phi(z) = 1 - \sum_{j=1}^{p*} (a_j + b_j) z^j$.
- Then the best predictor of X_t^2 is linear:

$$\mathbb{E}[X_t^2 | X_{t-1}] = \frac{a_0}{\theta(1)} + \sum_{k=1}^{\infty} \pi_k X_{t-k}^2,$$

where $\pi(z) = \phi(z)/\theta(z) = 1 - \sum_{k=1}^{\infty} \pi_k z^k$.

Remark 11.4.6. Fitting a GARCH Model

- Note that $Z_t = X_t/\sigma_t$, which is i.i.d. with some given probability density function p_Z .
- From a sample, we have realization x_t and compute

$$s_t^2 = \frac{a_0}{1 - b_1} + a_1 \sum_{j=0}^{\infty} b_1^j x_{t-1-j}^2,$$

truncating the sum where the sample begins.

- Then the *pseudo-likelihood* is defined as

$$\sum_{t=2}^n (\log p_Z(x_t/s_t) - \log s_t),$$

where we truncate the s_t calculation as needed.

- This pseudo-likelihood is a function of a_0, a_1, b_1 , and can be maximized.

Example 11.4.7. Fitting a Gaussian GARCH(1,1) to Dow Log Returns

- We consider the Dow Log Returns, and fit a Gaussian GARCH(1,1) model.

```
dow <- read.table("dow.dat")
dow <- diff(log(dow[,1]))
dow <- ts(dow, start=c(2008,164), frequency=252)
```

- Then we utilize the pseudo-likelihood based on a Gaussian marginal distribution.
- This uses a reparameterization of the GARCH parameters, to ensure we get values in $[0, 1]$ that satisfy the constraint that $\delta < 1$.

```
polymul <- function(a,b)
{
  bb <- c(b,rep(0,length(a)-1))
  B <- toeplitz(bb)
  B[lower.tri(B)] <- 0
  aa <- rev(c(a,rep(0,length(b)-1)))
  prod <- B %*% matrix(aa,length(aa),1)
  return(rev(prod[,1]))
}

psi2arch <- function(psi)
{
```



```

p <- length(psi)-1
a.0 <- exp(psi[1])
if(p > 0)
{
  r <- (1 + exp(-psi[2]))^(-1)
  if(p > 1)
  {
    a.1 <- (1 + sum(exp(-psi[3:(p+1)])))^(-1)
    a.j <- a.1
    for(j in 2:p)
    {
      a.j <- c(a.j,exp(-psi[j+1])*a.1)
    }
    a.j <- r*a.j
  } else { a.j <- r }
  a.0 <- c(a.0,a.j)
}
return(a.0)
}

lik.garch <- function(psi,data,p.order,df,fitting=TRUE)
{

  p <- p.order
  q <- length(psi)-p-1
  T <- length(data)
  coef <- psi2arch(psi)
  acoef <- coef[1:(p+1)]
  bcoef <- NULL
  if(q > 0) { bcoef <- coef[(p+2):(p+q+1)] }
  theta <- bcoef
  phi <- acoef[-1]
  psi <- c(1,ARMAtoMA(ar=theta,ma=NULL,lag.max=T))
  psi <- polymul(psi,phi)

  lik <- 0
  resids <- NULL
  for(t in (p+1):T)
  {
    new.sigt <- sqrt(acoef[1]/(1-sum(theta)) + sum(psi[1:(t-1)]*data[(t-1):1]^2))
    if(df==Inf) { lik <- lik + (-2)*log(dnorm(data[t],sd=new.sigt)) } else {
      lik <- lik + (-2)*log(dt(data[t]/new.sigt,df=df)/new.sigt) }
    resids <- c(resids,data[t]/new.sigt)
  }
  if(fitting) { return(lik) } else { return(resids) }
}

```

- We initialize and run the optimization.

```

p <- 1
q <- 1
psi.init <- rep(0,p+q+1)
garch.fit <- optim(psi.init,fn=lik.garch,data=dow,p.order=p,df=Inf,method="BFGS")
print(garch.fit)

```

```
## $par
## [1] -60.497149 13.003395 -2.306336
##
## $value
## [1] -12716.67
##
## $counts
## function gradient
##      48      22
##
## $convergence
## [1] 0
##
## $message
## NULL
```

```
garch.par <- psi2arch(garch.fit$par)
print(round(garch.par,digits=3))
```

```
## [1] 0.000 0.091 0.909
```

- The estimated parameters are: a_0 is $5.3262535 \times 10^{-27}$, a_1 is 0.0905994, and b_1 is 0.9093984.
- There is a high degree of persistence in volatility due to the high value of b_1 .

```
z <- lik.garch(garch.fit$par,dow,p,Inf,FALSE)
m <- length(z)
kurt <- m*sum((z - mean(z))^4)/(sum((z - mean(z))^2))^2 - 3
```

- Also $a_1 + b_1 \approx 1$, indicating high variance. We can check the residuals' kurtosis: it is 1.3230191. This high value indicates that the residuals have fat tails.

Example 11.4.8. Fat-Tailed GARCH(1,1) for Dow Log Returns

- The above fitted model assumes a Gaussian input, but the variance was large.
- So we should try a fat-tailed distribution. Consider Student t.

```
lik.tgarch <- function(psi,data,p.order)
{
  p <- p.order
  q <- length(psi)-p-2
  T <- length(data)
  df <- 2 + exp(psi[p+q+2])
  coef <- psi2arch(psi[1:(p+q+1)])
  acoef <- coef[1:(p+1)]
  bcoef <- NULL
  if(q > 0) { bcoef <- coef[(p+2):(p+q+1)] }
  theta <- bcoef
  phi <- acoef[-1]
  psi <- c(1,ARMAtoMA(ar=theta,ma=NULL,lag.max=T))
  psi <- polymul(psi,phi)

  lik <- 0
  for(t in (p+1):T)
  {
    new.sigt <- sqrt(acoef[1]/(1-sum(theta)) + sum(psi[1:(t-1)]*data[(t-1):1]^2))
    new.sigt <- new.sigt/sqrt(df/(df-2))
  }
}
```

```

    lik <- lik + (-2)*log(dt(data[t]/new.sig,df=df)/new.sig)
  }
  return(lik)
}

```

- We initialize and run the optimization.

```

p <- 1
q <- 1
# initialize with values close to those of the Gaussian GARCH
psi.init <- c(garch.fit$par,1)
tgarch.fit <- optim(psi.init,fn=lik.tgarch,data=dow,p.order=p,method="BFGS")
print(tgarch.fit)

## $par
## [1] -60.497149 13.020642 -2.167761 1.630332
##
## $value
## [1] -12816.1
##
## $counts
## function gradient
##      23      5
##
## $convergence
## [1] 0
##
## $message
## NULL

tgarch.par <- c(psi2arch(tgarch.fit$par[1:(p+q+1)]),2+exp(tgarch.fit$par[p+q+2]))
print(round(tgarch.par,digits=3))

## [1] 0.000 0.103 0.897 7.106

```

- The estimated parameters are: a_0 is $5.3262535 \times 10^{-27}$, a_1 is 0.102683, and b_1 is 0.8973148. The degrees of freedom is 7.1055719.
- We still have $a_1 + b_1 \approx 1$, indicating high variance.

Lesson 11-5: Bi-spectral Density

- Any Gaussian process is linear, and the 1st and 2nd cumulants (mean and covariance) summarize it completely.
- We can measure nonlinearity of a process through its 3rd cumulants, and the associated bi-spectral density.

Definition 11.5.1.

- Suppose $\{X_t\}$ is a strictly stationary time series with finite absolute third moment.
- Then its third *cumulant function* is defined to be

$$\gamma(k_1, k_2) = \mathbb{E}[(X_t - \mu)(X_{t+k_1} - \mu)(X_{t+k_2} - \mu)],$$

for any integer k_1, k_2 .

Remark 11.5.2. Cumulant Functions

- Cumulant functions generalize cumulants for random variables.
- The mean is the first cumulant function.
- The ACVF is the second cumulant function.
- The third cumulant function measures skewness across lags.

Fact 11.5.3. Symmetries of the Third Cumulant

- There are symmetries:

$$\gamma(k_1, k_2) = \gamma(k_2, k_1) = \gamma(k_1 - k_2, -k_2) = \gamma(-k_1, k_2 - k_1).$$

Definition 11.5.6.

The *bi-spectral density* function of a strictly stationary time series is the Fourier transform of the third cumulant function:

$$f(\lambda_1, \lambda_2) = \sum_{k_1, k_2 = -\infty}^{\infty} \gamma(k_1, k_2) \exp\{-i\lambda_1 k_1 - i\lambda_2 k_2\},$$

where $\lambda_1, \lambda_2 \in [-\pi, \pi]$.

Example 11.5.7. Bi-spectral Density of an ARCH(1) Process

- Consider the ARCH(1) process $X_t = Z_t \sqrt{a_0 + a_1 X_{t-1}^2}$.
- The bi-spectral density is

$$f(\lambda_1, \lambda_2) = \mu_3 \left((1 - a_1 e^{-i\lambda_1 - i\lambda_2})^{-1} + (1 - a_1 e^{i\lambda_1})^{-1} + (1 - a_1 e^{i\lambda_2})^{-1} - 2 \right),$$

where $\mu_3 = \mathbb{E}[X_t^3]$.

Theorem 11.5.8.

- Suppose that $\{X_t\}$ is a stationary linear time series such that $X_t = \mu + \psi(B)\epsilon_t$, with $\{\epsilon_t\}$ i.i.d. $(0, \sigma^2)$ and $\psi(z) = \sum_{j=-\infty}^{\infty} \psi_j z^j$.
- Then the bi-spectral density is

$$f(\lambda_1, \lambda_2) = \kappa_3 \psi(e^{-i\lambda_1}) \psi(e^{-i\lambda_2}) \psi(e^{i\lambda_1 + i\lambda_2}),$$

where $\kappa_3 = \mathbb{E}[\epsilon_t^3]$.

Lesson 11-6: Volatility Filtering

- ARCH and GARCH models do not *fully* capture volatility clustering and *fat tails* (i.e., extreme observations).
- We can use a model-free approach to describe σ_t , the volatility.
- In this lesson we view σ_t as deterministic, and defined via $\sigma_t = \mathbb{E}[X_t^2]$.

Definition 11.6.2

- Consider a moving average of squared returns:

$$\Pi(B)X_t^2 = \sum_{j=-\infty}^{\infty} \pi_j X_{t-j}^2.$$

This is a *volatility filter*, if all $\pi_j \geq 0$.

- A volatility filter can be causal, symmetric, etc.

Paradigm 11.6.3. Unbiased Volatility Filtering

- We consider estimating volatility $\sigma_t^2 = \mathbb{E}[X_t^2]$ by a volatility filter, just like we did for estimating the mean $\mu_t = \mathbb{E}[X_t]$:

$$\hat{\sigma}_t^2 = \Pi(B)X_t^2.$$


- The mean is $\Pi(B)\sigma_t^2$, which can be compared to the target σ_t^2 ; if it differs, there is bias.
- The same bias-variance trade-off applies as with mean estimation.

Example 11.6.4. Volatility Filtering of Dow Jones Log Returns

- Consider the volatility filter $\pi_j = 1/11$ for $|j| \leq 5$.
- Observe the increased volatility in more recent years.

```
dow <- read.table("dow.dat")
dow <- diff(log(dow[,1]))
dow <- ts(dow,start=c(2008,164),frequency=252)

p <- 5
dow.vol <- filter(dow^2,rep(1,2*p+1)/(2*p+1),method="convolution",sides=2)
plot(dow.vol,xlab="Year",ylab="",lwd=1)
```



Definition 11.6.5.

- When $\Pi(B)$ is causal, then $\hat{\sigma}_t^2$ depends only on present and past data.
- Then we can divide out by the (square root) volatility, getting a *pseudo-residual*:

$$W_t = \frac{X_t}{\hat{\sigma}_t}.$$

This is called the *Normalizing and Variance Stabilizing* (NoVaS) transform.

- The NoVaS denominator $\hat{\sigma}_t$ is unbiased if $\Pi(1) = 1$, i.e., the sum of all coefficients equals one.


Paradigm 11.6.6. NoVaS with Simple Weights

- The idea is to choose a NoVaS transform $\Pi(B)$ such that $\{W_t\}$ is close to i.i.d. Gaussian.
- We want light tails and no serial correlation.
- We can search over volatility filters to find pseudo-residuals with this property.
- NoVaS with simple weights: take $\Pi(B)$ causal with $\pi_0 = \pi_1 = \dots = \pi_p = 1/(p+1)$.
- We just need to choose p , which is done by making the kurtosis low (close to that of a Gaussian).

Example 11.6.7. Simple NoVaS Applied to Dow Log Returns

- Test out Simple NoVaS on the Dow log returns time series.

```
p <- 9
psi <- rep(1,p+1)/(p+1)
n <- length(dow)
dow.vol <- filter(dow^2,psi,method="convolution",sides=1)[(p+1):n]
z <- dow[(p+1):n]/sqrt(dow.vol)
m <- length(z)
plot(ts(z))
```




- It turns out, the serial correlation is low and the kurtosis is close to three (the Gaussian case)!
- We print out the excess kurtosis and a standardized excess kurtosis.

```
kurt <- m*sum((z - mean(z))^4)/(sum((z - mean(z))^2))^2 - 3
kurt.std <- sqrt(m)*kurt/sqrt(mean(z^8)-mean(z^4)^2)
```

```
z.acf <- acf(z,plot=TRUE,lag.max=m-1,type="correlation")$acf
```

Series z



```
print(c(kurt,kurt.std))
```

```
## [1] -0.2743338 -1.8086432
```