

Time Series: A First Course with Bootstrap Starter

Contents

Lesson 12-1: Sampling Distributions	2
Remark 12.1.4. Improving on the Normal Approximation	2
Example 12.1.6. Mean of a Gaussian AR(1).	2
Lesson 12-2: Monte Carlo	3
Fact 12.2.2. Parameters are Functionals of the Distribution	3
Example 12.2.5. Median Parameter	3
Example 12.2.12. Monte Carlo Approximation to the Variance of a Statistic	3
Exercise 12.6. Monte Carlo Approximation to the Variance of the Median	4
Lesson 13-3: The Plug-in Principle and the Bootstrap	5
Definition 12.3.1.	5
Paradigm 12.3.5. The Plug-In Principle	5
Example 12.3.8. Plug-In Estimator of the Median	5
Paradigm 12.3.9. Classical Bootstrap for the Variance of a Statistic	5
Example 12.3.10. Bootstrap for the Variance of U.S. Population Acceleration	6
Lesson 12-4: Model-based Bootstrap	7
Paradigm 12.4.2. Bootstrapping an AR(1) Model	7
Example 12.4.3. Bootstrap for the AR(1) Coefficient of U.S. Population Growth	7
Paradigm 12.4.7. Bootstrap and the Model-Free Principle	9
Lesson 12-5: Sieve Bootstraps	10
Paradigm 12.5.2. Sieves	10
Paradigm 12.5.4. Autoregressive Sieve and the AR Sieve Bootstrap	10
Example 12.5.6. Lag 12 Autocorrelation of Gasoline Sales	10
Paradigm 12.5.8. Linear Process Bootstrap	12
Example 12.5.11. Lag 1 Autocovariance of Non-Defense Capitalization	13
Lesson 12-6: Time Frequency Toggle Bootstrap	14
Paradigm 12.6.1. Spectral Sieve	14
Example 12.6.3. Wolfer Sunspot Spectral Sieve	14
Paradigm 12.6.8. Time Frequency Toggle Bootstrap	16
Example 12.6.11. Lag 12 Autocorrelation of Gasoline Sales via TFT	16
Lesson 12-7: Subsampling	18
Paradigm 12.7.1. Roots and Subsampling	18
Paradigm 12.7.2. Blocking Schemes	19
Paradigm 12.7.5. Subsampling Methodology	19
Example 12.7.13. Subsampling Inference for Lag 1 Autocovariance of Non-Defense Capitalization	19
Lesson 12-8: Block Bootstrap	21
Paradigm 12.8.2. Block Bootstrap for the Sample Mean	21
Exercise 12.33. Block Bootstrap for the Mauna Loa Mean	22

Lesson 12-1: Sampling Distributions

- In time series (and statistics in general) we want to know the uncertainty in our estimates.
- First we study the sampling distributions of statistics.

Remark 12.1.4. Improving on the Normal Approximation

- A confidence interval for the mean requires us to know the *long-run variance*,

$$\sigma_\infty^2 = \sum_{h=-\infty}^{\infty} \gamma(h).$$

- The 95% confidence interval based on the sample mean and the normal approximation is

$$\bar{X}_n \pm 1.96 \sqrt{\frac{\widehat{\sigma_\infty^2}}{n}}.$$

- There are some issues with the accuracy of this interval:
 1. The accuracy of the normal approximation in finite sample (non-normality of marginal distribution).
 2. Serial correlation in the time series, since

$$\text{Var}[\sqrt{n}\bar{X}_n] \rightarrow \sigma_\infty^2,$$

and the convergence can be slow.

3. Estimation of σ_∞^2 by some estimator $\widehat{\sigma_\infty^2}$ (this could be constructed from sample autocovariances).
- So we might consider directly approximating the distribution of \bar{X}_n via *resampling* from $\{X_t\}$. This is the idea of the *bootstrap*.

Example 12.1.6. Mean of a Gaussian AR(1).


- Suppose that $\{X_t\}$ is a stationary Gaussian AR(1) process, with parameter ϕ_1 .
- So the spectral density is $f(\lambda) = \sigma^2 |1 - \phi_1 e^{-i\lambda}|^{-2}$.
- The asymptotic variance of $\sqrt{n}\bar{X}$ is the long-run variance: $f(0) = \sigma^2/(1 - \phi_1)^2$.
- The actual variance of $\sqrt{n}\bar{X}$ is

$$\gamma(0) + 2 \sum_{h=1}^{n-1} (1 - h/n) \gamma(h) = \frac{\sigma^2}{1 - \phi_1^2} \left(1 + \frac{2\phi_1}{1 - \phi_1} \left[(1 - \phi_1^n) - \frac{1 - \phi_1^n(1 + n(1 - \phi_1))}{n(1 - \phi_1)} \right] \right)$$

- For $\phi_1 > 0$, this quantity is smaller than the long-run variance.
- We illustrate for $\phi_1 = .9$ and $\sigma = 1$, so that $f(0) = 100$.

```
phi <- .9
n <- seq(10,1000,10)

longrun.var <- (1-phi^2)^(-1)*(1 + 2*phi*(1-phi)^(-1)*
  ((1-phi^n) - (1-phi^n*(1+n*(1-phi)))*(n*(1-phi))^(-1)))
longrun.var <- ts(longrun.var,start=10,frequency=.10)
plot(longrun.var,xlab="Sample Size",ylab="",lwd=2,col=grey(.2))
abline(h=1/(1-phi)^2,lty=3,lwd=2)
```



- Thus, using $f(0)$ instead of the true sampling variance will overestimate when n is small.

Lesson 12-2: Monte Carlo

- For this lesson, we suppose that X_1, X_2, \dots, X_n are i.i.d. with common distribution G .

Fact 12.2.2. Parameters are Functionals of the Distribution

- Any parameter θ of a distribution can be expressed in terms of that distribution.
- Let G be the cumulative distribution function (cdf). Then we write θ as $\theta(G)$.

Example 12.2.5. Median Parameter

The median θ can be written as $G^{-1}(1/2)$, where G^{-1} denotes the quantile inverse. ($G^{-1}(p) = \inf\{x : G(x) \geq p\}$.)

Example 12.2.12. Monte Carlo Approximation to the Variance of a Statistic

- Suppose we want to know the variance of a statistic, $\eta = \text{Var}[\hat{\theta}_n]$, but there is no nice analytic formula.
- Monte Carlo approach: generate multiple independent copies of $\hat{\theta}_n$, and take the sample variance of these:

1. For a large integer M simulate:

$$\begin{aligned} X_1^{(1)}, X_2^{(1)}, \dots, X_n^{(1)} &\sim \text{i.i.d. } G \\ X_1^{(2)}, X_2^{(2)}, \dots, X_n^{(2)} &\sim \text{i.i.d. } G \\ &\vdots \\ X_1^{(M)}, X_2^{(M)}, \dots, X_n^{(M)} &\sim \text{i.i.d. } G. \end{aligned}$$

2. For $j = 1, \dots, M$ compute $\hat{\theta}_n^{(j)}$ from the pseudo-data $X_1^{(j)}, X_2^{(j)}, \dots, X_n^{(j)}$.
3. Compute $\hat{\mathbb{E}}[\hat{\theta}_n] = M^{-1} \sum_{j=1}^M \hat{\theta}_n^{(j)}$ and $\hat{\eta} = M^{-1} \sum_{j=1}^M (\hat{\theta}_n^{(j)} - \hat{\mathbb{E}}[\hat{\theta}_n])^2$, which is our estimate of η .

Exercise 12.6. Monte Carlo Approximation to the Variance of the Median

- We use Monte Carlo to approximate the variance of the sample median.
- Consider a sample of size $n = 100$ from an AR(1) process with mean 2, AR parameter $\phi_1 = .8$, and Student t inputs with 4 degrees of freedom.
- First we load the function to simulate an ARMA with Student t inputs.

```
armapq.simht <- function(n,burn,ar.coefs,ma.coefs,innovar,df,seed)
{
  p <- length(ar.coefs)
  q <- length(ma.coefs)
  set.seed(seed)
  if(df == Inf)
  {
    z <- rnorm(n+burn+p+q,sd=sqrt(innovar))
  } else
  {
    z <- sqrt(innovar)*rt(n+burn+p+q,df=df)
  }
  x <- filter(z,c(1,ma.coefs),method="convolution",sides=1)
  x <- x[(q+1):(q+n+burn+p)]
  y <- x[1:p]
  for(t in (p+1):(p+n+burn))
  {
    next.y <- sum(ar.coefs*y[(t-1):(t-p)]) + x[t]
    y <- c(y,next.y)
  }
  y <- y[(p+burn+1):(p+burn+n)]
  return(y)
}

n <- 100
phi1 <- .8
theta <- 2
monte <- 10000
med.mcs <- NULL
for(i in 1:monte)
{
  x.sim <- theta + armapq.simht(n,500,phi1,NULL,1,4,set.seed(i))
  med.mcs <- c(med.mcs,median(x.sim))
}
print(mean(med.mcs^2)-(mean(med.mcs))^2)
```

Lesson 13-3: The Plug-in Principle and the Bootstrap

- We describe the plug-in method of constructing estimators, which generalizes the method-of-moments.

Definition 12.3.1.

- If X_1, \dots, X_n are i.i.d. with common cdf G , then their empirical distribution function (edf) is

$$\hat{G}(x) = \frac{1}{n} \sum_{i=1}^n 1_{\{X_i \leq x\}}.$$

Paradigm 12.3.5. The Plug-In Principle

- The edf converges uniformly to the cdf, so we can estimate θ by *plugging in*:

$$\hat{\theta} = \theta(\hat{G}).$$

- This is called a *plug-in estimator*.
- It generalizes method-of-moments.

Example 12.3.8. Plug-In Estimator of the Median

The plug-in estimator for the median is

$$\hat{\theta} = \hat{G}^{-1}(1/2),$$

which is the sample median.

Paradigm 12.3.9. Classical Bootstrap for the Variance of a Statistic

- Consider a scenario where we want to know the variance of a statistic $\hat{\theta}_n$ (computed from a sample of size n).
- Let $\eta = \text{Var}[\hat{\theta}_n]$. Since n is fixed and the sample is i.i.d., $\eta = \eta(G)$.
- We could estimate η with the plug-in estimator.
- We would like to compute $\eta(\hat{G})$, but maybe there is no formula!
- So we try to approximate it using the *bootstrap*.

1. For large M simulate

$$\begin{aligned} X_1^{*(1)}, \dots, X_n^{*(1)} &\sim \text{i.i.d.} \hat{G} \\ X_1^{*(2)}, \dots, X_n^{*(2)} &\sim \text{i.i.d.} \hat{G} \\ &\vdots \\ X_1^{*(M)}, \dots, X_n^{*(M)} &\sim \text{i.i.d.} \hat{G}. \end{aligned}$$


2. For $1 \leq j \leq M$ compute $\hat{\theta}_n^{*(j)}$ from the *pseudo-sample* $X_1^{*(j)}, \dots, X_n^{*(j)}$.
3. Our bootstrap estimator of η is

$$\frac{1}{M} \sum_{j=1}^M \left(\hat{\theta}_n^{*(j)} - M^{-1} \sum_{k=1}^M \hat{\theta}_n^{*(k)} \right)^2.$$

Example 12.3.10. Bootstrap for the Variance of U.S. Population Acceleration

- Consider the time series $\{Y_t\}$ of U.S. Population.
- One possible model is twice differencing, for which the series appears to be white noise.

```
pop <- read.table("USpop.dat")
pop <- ts(pop, start = 1901)
pop.diff <- diff(diff(pop))*10^(-3)
acf(pop.diff)
```



- Suppose that $X_t = (1 - B)^2 Y_t$ is actually i.i.d.
- The sample mean of $\{X_t\}$ is 9.8227835 in units of millions.
- Suppose we want to estimate the variance of this sample mean using the bootstrap (of course we could use a formula as well) with $M = 10^5$.

```
n <- length(pop.diff)
pop.mean <- mean(pop.diff)
pop.edf <- sort(pop.diff)

monte.means <- NULL
Monte <- 100000
for(i in 1:Monte)
{
  monte.sample <- sample(pop.edf,size=n,replace=TRUE)
  monte.means <- c(monte.means,mean(monte.sample))
}
var.mean <- var(monte.means)
```

- The resulting estimate is 627.3034894 in units of millions.
- Note: results change each time notebook is rendered, because seed is not fixed!

Lesson 12-4: Model-based Bootstrap

- We want to extend the bootstrap idea to the case of time series data.
- We don't want to assume i.i.d. anymore, because time series data have serial dependence.

Paradigm 12.4.2. Bootstrapping an AR(1) Model

- Suppose $\{X_t\}$ is a stationary AR(1) process with i.i.d. inputs with cdf G :

$$X_t - \phi X_{t-1} = \epsilon_t,$$

where $\epsilon_t \sim \text{i.i.d.} G$.

- Suppose we want to estimate the cdf of $\hat{\phi} - \phi$, where $\hat{\phi}$ is the Yule-Walker estimator.
- So for any x , we want a bootstrap estimate of $\zeta = \mathbb{P}[\hat{\phi} - \phi \leq x]$.
- We compute residuals

$$e_t = X_t - \hat{\phi} X_{t-1},$$

and use these as proxies for ϵ_t .

- Center the residuals by their sample mean, and let \hat{G} be their edf.
1. For large M simulate

$$\begin{aligned} \epsilon_1^{*(1)}, \dots, \epsilon_n^{*(1)} &\sim \text{i.i.d.} \hat{G} \\ \epsilon_1^{*(2)}, \dots, \epsilon_n^{*(2)} &\sim \text{i.i.d.} \hat{G} \\ &\dots \\ \epsilon_1^{*(M)}, \dots, \epsilon_n^{*(M)} &\sim \text{i.i.d.} \hat{G}. \end{aligned}$$

2. For $1 \leq j \leq M$ construct

$$X_t^{*(j)} = \hat{\phi} X_{t-1}^{*(j)} + \epsilon_t^{*(j)}$$

for $1 \leq t \leq n$.

3. Compute $\hat{\phi}^{*(j)}$ from the *pseudo-sample* $X_1^{*(j)}, \dots, X_n^{*(j)}$.
4. Our bootstrap estimator of ζ is


$$\frac{1}{M} \sum_{j=1}^M 1_{\{\hat{\phi}^{*(j)} - \hat{\phi} \leq x\}}.$$

Example 12.4.3. Bootstrap for the AR(1) Coefficient of U.S. Population Growth

- For the U.S. Population time series $\{Y_t\}$, we can also consider fitting an AR(1) model to first differences.

```
pop <- read.table("USpop.dat")
pop <- ts(pop, start = 1901)
pop.diff <- diff(pop)*10^(-6)
acf(pop.diff)
```

V1




- Let $X_t = (1 - B)Y_t$, and consider the AR(1) model for $\{X_t\}$.
- Suppose we want to estimate the cdf of $\hat{\phi} - \phi$ using the bootstrap with $M = 10^5$.

```
n <- length(pop.diff)
kappa.hat <- pacf(pop.diff,lag=n-1,plot=FALSE)$acf[,1]
pop.ar1 <- kappa.hat[1]
pop.resids <- pop.diff[2:n] - pop.ar1*pop.diff[1:(n-1)]
pop.resids <- pop.resids - mean(pop.resids)
pop.edf <- sort(pop.resids)

monte.roots <- NULL
Monte <- 100000
for(i in 1:Monte)
{
  monte.resids <- sample(pop.edf,size=n,replace=TRUE)
  init.value <- sample(pop.diff,size=1)
  monte.sample <- filter(monte.resids,pop.ar1,method="recursive",init=init.value)
  monte.root <- pacf(monte.sample,lag=n-1,plot=FALSE)$acf[,1][1] - pop.ar1
  monte.roots <- c(monte.roots,monte.root)
}
# hist(monte.roots)
interval <- c(sort(monte.roots)[floor(.025*Monte)],sort(monte.roots)[floor(.975*Monte)])
```

- The AR(1) coefficient is estimated to be 0.9130219.
- The 95% confidence interval based on the bootstrap is [0.8838385,1.0534122].
- We plot the bootstrap edf.


```
plot(sort(monte.roots),seq(1,Monte)/Monte,type="l",xlab="x",ylab="",lwd=2)
```



Paradigm 12.4.7. Bootstrap and the Model-Free Principle

- Using a transformation (instead of a model) that produces i.i.d. residuals from the data process is called the *model-free principle*.
- So to do a time series bootstrap, we should seek such a transformation, bootstrap the residuals, reconstruct the process, and evaluate the statistic on the pseudo-samples.
- Suppose there exists an invertible transformation Π such that $\underline{\epsilon} = \Pi(\underline{X})$ is a vector of i.i.d. components, where $\underline{X} = [X_1, \dots, X_n]'$.
- Let G denote the cdf of ϵ_t .
- Suppose we have a statistic $\hat{\theta}_n$ and we want the cdf $\zeta = \mathbb{P}[\hat{\theta}_n - \theta \leq x]$.
- Compute the residuals, and estimate G via the residual edf \hat{G} .
- Then the *Model-free bootstrap* is:

1. For large M simulate

$$\begin{aligned} \epsilon_1^{*(1)}, \dots, \epsilon_n^{*(1)} &\sim \text{i.i.d.} \hat{G} \\ \epsilon_1^{*(2)}, \dots, \epsilon_n^{*(2)} &\sim \text{i.i.d.} \hat{G} \\ &\dots \\ \epsilon_1^{*(M)}, \dots, \epsilon_n^{*(M)} &\sim \text{i.i.d.} \hat{G}. \end{aligned}$$

2. For $1 \leq j \leq M$ construct

$$\underline{X}^{*(j)} = \Pi^{-1}[\underline{\epsilon}^{*(j)}].$$

3. Compute $\hat{\theta}_n^{*(j)}$ from the *pseudo-sample* $\underline{X}^{*(j)}$.
4. Our bootstrap estimator of ζ is

$$\frac{1}{M} \sum_{j=1}^M 1_{\{\hat{\theta}_n^{*(j)} - \hat{\theta} \leq x\}}.$$

Lesson 12-5: Sieve Bootstraps

- We now investigate two transformations Π : the AR sieve and the MA sieve.

Paradigm 12.5.2. Sieves

- Consider the case that the transformation Π involves infinitely many parameters.
- So we consider a sequence of transformations Π_1, Π_2, \dots , where Π_j has j parameters.
- We suppose these transformations to be nested. This means that Π_j is obtained from Π_{j+1} by restricting the $j+1$ th parameter to some constant value (such as zero).
- Also we suppose that taking the limit of these transformations gives Π .
- Such a collection is called a *sieve*.
- The *method of sieves* is to apply Π_m to the sample X_1, \dots, X_n , with m chosen large enough that $\Pi_m \approx \Pi$, while also m is small enough that we can estimate all the parameters.
- If we get additional data (n increases), then we would also increase m .

Paradigm 12.5.4. Autoregressive Sieve and the AR Sieve Bootstrap

- Suppose we have an $\text{AR}(\infty)$ process:

$$\Xi(B)X_t = \epsilon_t \sim \text{i.i.d.} G.$$

- We take Π_p to be an $\text{AR}(p)$ model. This is called the *AR sieve*.
- Notation:

$$\phi^{(p)}(B) = 1 - \sum_{j=1}^p \phi_j^{(p)} B^j.$$

- So as p increases, all the coefficients can change (and we get more coefficients, too).
- The order p can be linked to sample size n by a formula. Or p can be determined empirically, as a statistic of the sample.
- Once this Π_p is determined, we can do an $\text{AR}(p)$ bootstrap (generalizing the $p = 1$ case considered in previous notebook).
- Because p grows with n , this is called the *AR sieve bootstrap*.

Example 12.5.6. Lag 12 Autocorrelation of Gasoline Sales

- Consider the seasonally adjusted gasoline sales data.
- Apply logs and differences, obtaining a linear process (this is an assumption).

```
gassa <- read.table("GasSA_2-11-13.dat")
gassa <- ts(log(gassa), start=1992, frequency=12)
gas.diff <- diff(gassa)
n <- length(gas.diff)
```

- We want to estimate $\rho(12)$, and get the cdf of $\hat{\rho}(12) - \rho(12)$.
- We use an AR sieve bootstrap with $p = 12$ (based on analysis of the PACF plot) for this sample size. There are $M = 10^5$ replications, and the pseudo-samples are constructed using a burnin of 500.

```

rho.hat <- acf(gas.diff,lag=n-1,type="correlation",plot=FALSE)$acf[,1]
gas.acf12 <- rho.hat[13]
p.order <- 12
phi.ar <- solve(toeplitz(rho.hat[1:p.order])) %*% rho.hat[2:(p.order+1)]

gas.resids <- gas.diff[(p.order+1):n]
for(i in 1:p.order) { gas.resids <- gas.resids - phi.ar[i]*gas.diff[(p.order+1-i):(n-i)] }
gas.resids <- gas.resids - mean(gas.resids)
gas.edf <- sort(gas.resids)

monte.roots <- NULL
burnin <- 500
Monte <- 100000
for(i in 1:Monte)
{
  monte.resids <- sample(gas.edf,size=n+burnin,replace=TRUE)
  init.value <- rep(0,p.order)
  monte.sample <- filter(monte.resids,phi.ar,method="recursive",init=init.value)[(burnin+1):(burnin+n)]
  monte.root <- acf(monte.sample,lag=n-1,plot=FALSE)$acf[,1][13] - gas.acf12
  monte.roots <- c(monte.roots,monte.root)
}
# hist(monte.roots)
interval <- c(sort(monte.roots)[floor(.025*Monte)],sort(monte.roots)[floor(.975*Monte)])


```

- The lag 12 autocorrelation is estimated to be -0.239469.
- The 95% confidence interval based on the bootstrap is [-0.3850592,-0.1226657].
- We plot the bootstrap edf.

```

plot(sort(monte.roots),seq(1,Monte)/Monte,type="l",xlab="x",ylab="",lwd=2)

```



Paradigm 12.5.8. Linear Process Bootstrap

- Consider a stationary process $\{Y_t\}$ with mean μ and acvf $\gamma(h)$.
- For a sample of size n , we have Γ_n is the Toeplitz covariance matrix of the sample. Recall that $\Gamma_n = \{\gamma(j - k)\}$.
- We can taper the sample autocovariance estimators in order to estimate the whole matrix:

$$\check{\gamma}(h) = \Lambda(h/d)\hat{\gamma}(h).$$

- Here Λ is a *taper*, which is a symmetric function on $[-1, 1]$ with non-negative values, which down-weights $\hat{\gamma}$ when $|h|$ is large.
- Also d is the *bandwidth*, which is chosen by the user, and typically satisfies $d/n \rightarrow 0$.
- For example, Λ can be a trapezoid function.
- Then we construct $\check{\Gamma}_n$ by inserting $\check{\gamma}(h)$ for $\gamma(h)$, and ensuring the matrix is positive-definite.
- There is a Cholesky decomposition of the matrix, of the form

$$\check{\Gamma} = L D L',$$

where L is unit lower-triangular and D is diagonal with positive entries.

- Then we can transform the data to residuals by first subtracting the sample mean, and then multiplying the sample vector by $D^{-1/2}L^{-1}$.
- To the resulting residuals we apply the i.i.d. bootstrap; this whole procedure is called the *Linear process bootstrap*.
- When using a trapezoidal taper, $\check{\gamma}(h) = 0$ if $|h| > d$, so we can think of these autocovariance estimates as corresponding to an $MA(d)$ process.
- If we use some taper such that $\check{\gamma}(h)$ truncates at $h = q$, the structure resembles that of an $MA(q)$ process, and the resulting procedure is called the *MA sieve bootstrap*.

Example 12.5.11. Lag 1 Autocovariance of Non-Defense Capitalization

- Consider the Non-Defense Capitalization time series.
- After differencing, we wish to estimate $\gamma(1)$ and compute the cdf $\zeta = \mathbb{P}[\hat{\gamma}(1) - \gamma(1) \leq x]$.

```
ndc <- read.table("Nondefcap.dat")
ndc <- ts(ndc[,2],start=c(1992,3),frequency=12,names= "NewOrders")
ndc.diff <- diff(ndc)
n <- length(ndc.diff)
```

- We use an MA sieve bootstrap with the truncation taper and $q = 10$ for this sample size. There are $M = 10^5$ replications.


```
gamma.hat <- acf(ndc.diff,lag=n-1,type="covariance",plot=FALSE)$acf[,1]
ndc.acf1 <- gamma.hat[2]
q.order <- 10
gamma.mat <- toeplitz(c(gamma.hat[1:(q.order+1)],rep(0,n-(q.order+1))))

gamma.chol <- t(chol(gamma.mat))
ndc.resids <- solve(gamma.chol,ndc.diff)
ndc.resids <- ndc.resids - mean(ndc.resids)
ndc.edf <- sort(ndc.resids)

monte.roots <- NULL
Monte <- 100000
for(i in 1:Monte)
{
  monte.resids <- sample(ndc.edf,size=n,replace=TRUE)
  monte.sample <- gamma.chol %*% monte.resids
  monte.root <- acf(monte.sample,lag=n-1,plot=FALSE,type="covariance")$acf[,1][2] - ndc.acf1
  monte.roots <- c(monte.roots,monte.root)
}
# hist(monte.roots)
interval <- c(sort(monte.roots)[floor(.025*Monte)],sort(monte.roots)[floor(.975*Monte)])
```

- The lag 1 autocovariance is estimated to be -0.002489.
- The 95% confidence interval based on the bootstrap is [-0.003506,-0.0010687].
- We plot the bootstrap edf.

```
plot(sort(monte.roots),seq(1,Monte)/Monte,type="l",xlab="x",ylab="",lwd=2)
```



Lesson 12-6: Time Frequency Toggle Bootstrap

- We study a frequency domain sieve.

Paradigm 12.6.1. Spectral Sieve

- The DFT $\tilde{\underline{X}}$ is obtained from the sample \underline{X} via multiplication by Q :

$$\tilde{\underline{X}} = Q \underline{X}.$$

- The DFT has asymptotic covariance given by diagonal matrix Λ ; the diagonal entries are the spectral density evaluated at Fourier frequencies.
- So $\Lambda^{-1/2} \tilde{\underline{X}}$ is asymptotically uncorrelated with common variance 1.
- Then

$$\underline{\epsilon} = Q \Lambda^{-1/2} Q^* \underline{X}$$

is an entropy-increasing transformation of the sample, and is an asymptotically uncorrelated sequence.

- We can base a bootstrap on $\underline{\epsilon}$; this is called a *spectral sieve*, because it involves an approximation to the true spectral density, which improves as the grid of Fourier frequencies becomes finer.
- To implement, we need an estimate of the spectral density.

Example 12.6.3. Wolfer Sunspot Spectral Sieve

- Consider the Wolfer sunspot time series.
- We estimate the DFT.

```

wolfer <- read.table("wolfer.dat")
wolfer <- ts(wolfer,start=1749,frequency=12)
n <- length(wolfer)
gamma.hat <- acf(wolfer,lag=n-1,type="covariance",plot=FALSE)$acf[,1]

lambda <- seq(-floor(n/2)+1,floor(n/2))*2*pi/n
Q.mat <- exp(1i*seq(-floor(n/2)+1,floor(n/2)) %x% t(lambda))/sqrt(n)
wolfer.dft <- Conj(Q.mat) %*% (wolfer-mean(wolfer))

```

- We apply a tapered spectral estimator with Bartlett taper to estimate f .

```

d <- 3*floor(n^{1/3})
wolfer.spec <- cos(0*lambda)*gamma.hat[1]
for(h in 1:(n-1))
{
  wolfer.spec <- wolfer.spec + 2*(max(1-h/d,0))*cos(h*lambda)*gamma.hat[h+1]
}
wolfer.spec <- ts(wolfer.spec,start=0,frequency=n)

```


- Then we construct the residuals, and check that they are a white noise.
- The decorrelation is imperfect, but is sufficiently low statistically.

```

wolfer.sieved <- wolfer.dft/sqrt(wolfer.spec)
wolfer.resids <- Re(Q.mat %*% wolfer.sieved)
gamma.resids <- acf(wolfer.resids,lag=n-1,type="correlation")$acf[,1]

```

V1



Paradigm 12.6.8. Time Frequency Toggle Bootstrap

- We construct a bootstrap based on the spectral sieve.
- Let \hat{f} be a spectral density estimate, and $\hat{\Lambda}$ is diagonal with entries given by evaluating at the Fourier frequencies.
- The *frequency domain residuals* are defined as $\hat{\Lambda}^{-1/2}\tilde{X}$; these are asymptotically i.i.d. and Gaussian, by theory from Chapter 9.
- So we can generate bootstrap copies of the frequency domain residuals, transform to time domain by applying Q , and evaluate our statistic.
- This is called the *time frequency toggle* (TFT) bootstrap.

Example 12.6.11. Lag 12 Autocorrelation of Gasoline Sales via TFT

- Consider the Gasoline sales time series data.
- We want to estimate the lag 12 autocorrelation, and measure the uncertainty using the TFT bootstrap.

```
gassa <- read.table("GasSA_2-11-13.dat")
gassa <- ts(log(gassa),start=1992,frequency=12)
gas.diff <- diff(gassa)
n <- length(gas.diff)
gamma.hat <- acf(gas.diff,lag=n-1,type="covariance",plot=FALSE)$acf[,1]
rho.hat <- acf(gas.diff,lag=n-1,type="correlation",plot=FALSE)$acf[,1]
gas.acf12 <- rho.hat[13]
print(gas.acf12)
```


```
## [1] -0.239469
```

- We compute the DFT, and construct a spectral estimator based on a fitted AR(12).

```
lambda <- seq(-floor(n/2),floor(n/2))*2*pi/n
Q.mat <- exp(1i*seq(-floor(n/2),floor(n/2))) %x% t(lambda))/sqrt(n)
gas.dft <- Conj(Q.mat) %*% (gas.diff-mean(gas.diff))

p.order <- 12
phi.ar <- solve(toeplitz(gamma.hat[1:p.order])) %*% gamma.hat[2:(p.order+1)]
sig2.ar <- gamma.hat[1] - sum(phi.ar*gamma.hat[2:(p.order+1)])
gas.spec <- rep(1,n)
for(j in 1:p.order) { gas.spec <- gas.spec - phi.ar[j]*exp(-1i*j*lambda) }
gas.spec <- sig2.ar*Mod(gas.spec)^{-2}
gas.sieved <- gas.dft/sqrt(gas.spec)
#plot(ts(Re(gas.sieved)))
#ceps.sieve <- Q.mat %*% Lambda.mat %*% Conj(Q.mat)
gas.resids <- Re(Q.mat %*% gas.sieved)
gas.tft <- gas.sieved[1:(floor(n/2)+1)]
gamma.resids <- acf(gas.resids,lag=n-1,type="correlation")$acf[,1]
```


V1




- Next, we do the bootstrapping on the frequency domain residuals.

```
gas.tft <- gas.tft - mean(gas.tft)
gas.edf.re <- sort(Re(gas.tft))
gas.edf.im <- sort(Im(gas.tft))

monte.roots <- NULL
Monte <- 100000
for(i in 1:Monte)
{
  monte.resids <- sample(gas.edf.re,size=(floor(n/2)+1),replace=TRUE) +
    1i*sample(gas.edf.im,size=(floor(n/2)+1),replace=TRUE)
  monte.resids[(floor(n/2)+1)] <- Re(monte.resids[(floor(n/2)+1)])
  monte.resids <- c(monte.resids,Conj(rev(monte.resids)[-1]))
  monte.sample <- Q.mat %*% monte.resids*sqrt(gas.spec)
  monte.sample <- Re(monte.sample)
  monte.root <- acf(monte.sample,lag=n-1,plot=FALSE,type="correlation")$acf[,1][13] - gas.acf12
  monte.roots <- c(monte.roots,monte.root)
}
print(c(sort(monte.roots)[floor(.025*Monte)],sort(monte.roots)[floor(.975*Monte)]))

## [1] 0.1176180 0.3520214

plot(sort(monte.roots),seq(1,Monte)/Monte,type="l",xlab="x",ylab="",lwd=2,col=grey(.2))
```



Lesson 12-7: Subsampling

- An alternative to the bootstrap is to replicate properties of the sample by examining *subsamples*.

Paradigm 12.7.1. Roots and Subsampling

- Suppose that $\{X_t\}$ is strictly stationary.
- We also suppose that the process is m -dependent, which says that serial dependence vanishes between variables that are more than lag m apart.
- Suppose $\hat{\theta}_n$ is an estimator of θ based on the sample X_1, \dots, X_n .
- Suppose that τ_n is a rate of convergence for the estimator, such that

$$\tau_n(\hat{\theta}_n - \theta) \Rightarrow S,$$

where S is some random variable with cdf J .

- The cdf of our centered statistic is

$$J_n(x) = \mathbb{P}[\tau_n(\hat{\theta}_n - \theta) \leq x].$$

- So $J_n(x) \rightarrow J(x)$ as $n \rightarrow \infty$.
- We can compute the statistic on smaller sub-spans of the full sample, and the corresponding cdf will also converge.

Paradigm 12.7.2. Blocking Schemes

- We can divide X_1, \dots, X_n into smaller blocks.
- Here we focus on overlapping blocks of size b , where $b < n$.
- There are $Q = n - b + 1$ such blocks, or *subsamples*. For $1 \leq i \leq Q$

$$X_i, \dots, X_{i+b-1}$$

is the i th subsample.

- Adjacent blocks have $b - 1$ values in common.

Paradigm 12.7.5. Subsampling Methodology

- Consider overlapping blocks, and evaluate the statistic on each:

$$\hat{\theta}_{b,i} = \hat{\theta}(X_i, \dots, X_{i+b-1}).$$

- The centered statistic is then

$$Z_{b,i} = \tau_b(\hat{\theta}_{b,i} - \theta).$$

Note the rate is τ_b , not τ_n .

- As $b \rightarrow \infty$, the cdf of $Z_{b,i}$ tends to J , for each i .
- Fixing b , each $Z_{b,i}$ has the same distribution, and they are dependent random variables. The idea is to take their edf to estimate $J_n(x)$.
- However, θ in $Z_{b,i}$ is unknown, so we replace it by $\hat{\theta}_n$ based on the whole sample, which converges at a faster rate.

$$\hat{Z}_{b,i} = \tau_b(\hat{\theta}_{b,i} - \hat{\theta}_n).$$

- The *classical subsampling estimator* is then defined as

$$L_{n,b}(x) = \frac{1}{Q} \sum_{i=1}^Q 1_{\{\hat{Z}_{b,i} \leq x\}}.$$

- This is an edf. Its quantiles are the order statistics of $\hat{Z}_{b,i}$.
- Neighboring blocks share lots of time series observations, and the corresponding $\hat{Z}_{b,i}$ random variables will be more highly correlated.
- By the m -dependence assumption, blocks that are sufficiently separated (far apart from one another) will be uncorrelated.
- Assuming that $\tau_b/\tau_n \rightarrow 0$, the result is a consistent estimator of $J(x)$, and the difference $L_{n,b}(x) - J_n(x)$ tends to zero in probability.

Example 12.7.13. Subsampling Inference for Lag 1 Autocovariance of Non-Defense Capitalization

- Consider the Non-Defense Capitalization time series.
- After differencing, we wish to estimate $\gamma(1)$ and compute the cdf $\zeta = \mathbb{P}[\hat{\gamma}(1) - \gamma(1) \leq x]$.

```
ndc <- read.table("Nondefcap.dat")
ndc <- ts(ndc[,2],start=c(1992,3),frequency=12,names="NewOrders")
ndc.diff <- diff(ndc)
n <- length(ndc.diff)
gamma.hat <- acf(ndc.diff,lag=n-1,type="covariance",plot=FALSE)$acf[,1]
ndc.acf1 <- gamma.hat[2]
```

- The lag 1 autocovariance is estimated to be -0.002489.
- We use the subsampling methodology, with $\tau_n = \sqrt{n}$.
- First consider $b = 5$.

```

b.sub <- 5
q.sub <- floor(n-b.sub+1)
sub.edf <- NULL
for(i in 1:q.sub)
{
  sub.ndc <- ndc.diff[i:(i+b.sub-1)]
  gamma.sub <- acf(sub.ndc,lag=b.sub-1,type="covariance",plot=FALSE)$acf[,1]
  sub.edf <- c(sub.edf,sqrt(b.sub)*(gamma.sub[2]-ndc.acf1))
}
sub.edf <- sort(sub.edf)
interval <- c(sub.edf[floor(.025*q.sub)],sub.edf[floor(.975*q.sub)])


```

- The 95% confidence interval based on subsampling is $[-0.0089642, 0.0133861]$.
- We plot the subsampling edf.

```

plot(sub.edf,seq(1,q.sub)/q.sub,type="l",xlab="x",ylab="",lwd=2)

```



- Repeat with $b = 10$.

```

b.sub <- 10
q.sub <- floor(n-b.sub+1)
sub.edf <- NULL
for(i in 1:q.sub)
{
  sub.ndc <- ndc.diff[i:(i+b.sub-1)]
  gamma.sub <- acf(sub.ndc,lag=b.sub-1,type="covariance",plot=FALSE)$acf[,1]
  sub.edf <- c(sub.edf,sqrt(b.sub)*(gamma.sub[2]-ndc.acf1))
}

```

```

}
sub.edf <- sort(sub.edf)
interval <- c(sub.edf[floor(.025*q.sub)], sub.edf[floor(.975*q.sub)])


```

- The 95% confidence interval based on subsampling is $[-0.0111949, 0.0567725]$.
- We plot the subsampling edf.

```

plot(sub.edf, seq(1, q.sub)/q.sub, type="l", xlab="x", ylab="", lwd=2)

```



Lesson 12-8: Block Bootstrap

- We now look at *block bootstrap methods*, which is similar in spirit to subsampling.
- We take sub-spans of the time series, and then patch them together to generate a synthetic time series that functions as our pseudo-sample.

Paradigm 12.8.2. Block Bootstrap for the Sample Mean

- Suppose we have a stationary time series $\{X_t\}$.
- Suppose θ is the mean, and $\hat{\theta}$ is the sample mean. Then

$$J_n(x) = \mathbb{P}[\sqrt{n}(\hat{\theta} - \theta) \leq x] \rightarrow J(x),$$

where $J(x) = \Phi(x/\sigma_\infty)$.

- Let b be small relative to n , as in the subsampling method.
- Divide the sample into $Q = n - b + 1$ overlapping blocks.
- We want to construct a length n pseudo-sample from length b blocks, so we require $k = n/b$ such blocks.

- Idea: randomly draw from the Q blocks, and then paste them together!
 - Denote those Q blocks of random variables as B_1, \dots, B_Q .
1. For $1 \leq j \leq M$, draw $B_1^{*(j)}, \dots, B_k^{*(j)}$ randomly (with replacement) from the set of available blocks B_1, \dots, B_Q .
 2. Concatenate for each j those draws, obtaining the j th pseudo-series $X_1^{*(j)}, \dots, X_n^{*(j)}$.
 3. Compute the statistic on each pseudo-series, denoted $\hat{\theta}_n^{*(j)}$.
 4. We estimate $J_n(x)$ with

$$\frac{1}{M} \sum_{j=1}^M 1_{\{\sqrt{n}(\hat{\theta}_n^{*(j)} - \hat{\theta}_n) \leq x\}}.$$

- Just like subsampling, we insert $\hat{\theta}_n$ for θ .
- There is a variant called the *tapered block bootstrap*, where observations at the edge of a block B_k get down-weighted.

Exercise 12.33. Block Bootstrap for the Mauna Loa Mean

- Consider annual differences of the logged Mauna Loa CO2 time series.
- We want to estimate the mean using block bootstrap. So we apply the above method to obtain the cdf of the centered statistic.

```

mau <- read.table("mauna.dat",header=TRUE,sep="")
mau <- ts(mau,start=1958,frequency=12)
mau.gr <- 1000*diff(log(mau),lag=12)
n <- length(mau.gr)
mau.mean <- mean(mau.gr)

```

- The sample mean (after rescaling) is 3.6382049.
- First try $b = \sqrt{n}$ and $M = 10^5$.
- We can also try $b = n^{1/3}$ and $b = n^{2/3}$ by rerunning the code chunk.

```

# Make choice of b, and then run
b <- ceiling(sqrt(n))
#b <- ceiling(n^(1/3))
#b <- ceiling(n^(2/3))


k <- ceiling(n/b)
Q.big <- n - b + 1
monte.roots <- NULL
Monte <- 100000
for(i in 1:Monte)
{
  monte.sample <- NULL
  for(j in 1:k)
  {
    boot.index <- sample(Q.big,size=1)
    monte.sample <- c(monte.sample,mau.gr[seq(boot.index,boot.index+b-1)])
  }
  monte.sample <- monte.sample[1:n]
  monte.root <- mean(monte.sample) - mau.mean
  monte.roots <- c(monte.roots,monte.root)
}
interval <- c(sort(monte.roots)[floor(.025*Monte)],sort(monte.roots)[floor(.975*Monte)])

```

- The 95% confidence interval based on the bootstrap is [3.0408972,4.1694835].

- We plot the bootstrap edf.

```
plot(sort(monte.roots),seq(1,Monte)/Monte,type="l",xlab="x",ylab="",lwd=2)
```






Figure 1: Wrapping a time series around the circle